



# Workshop online Series

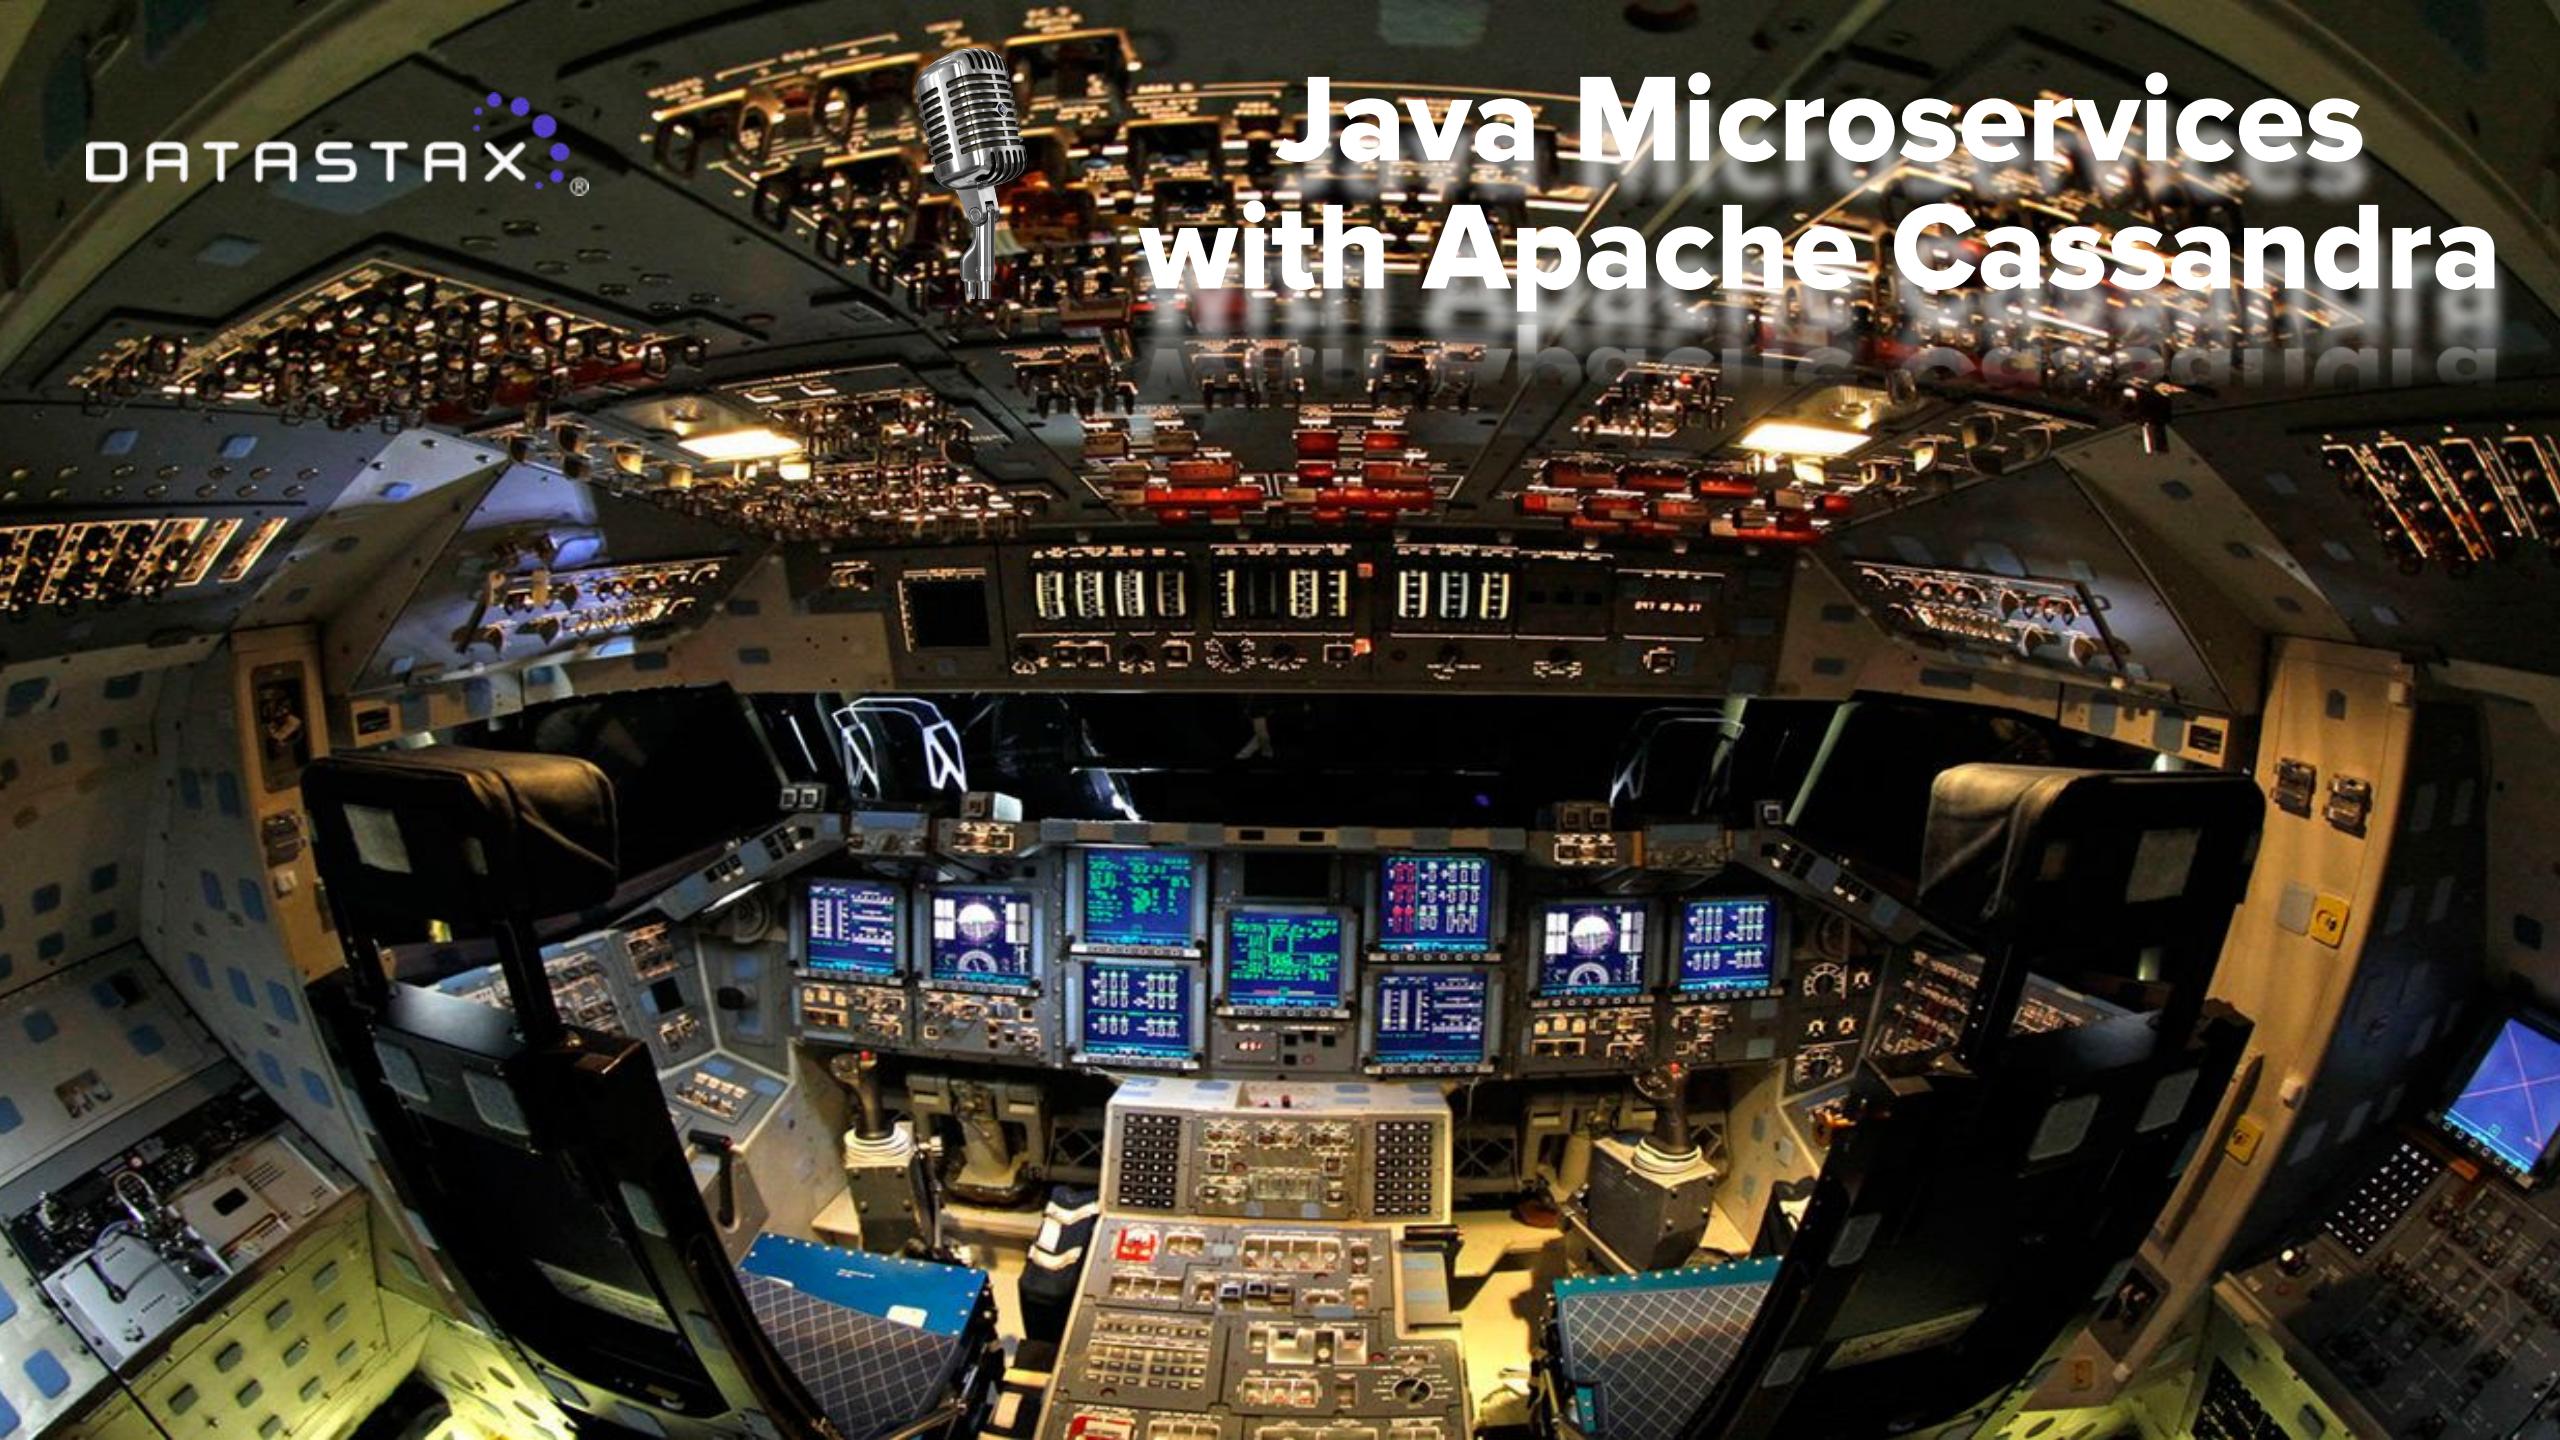


## Build Java Microservices for Apache Cassandra™





# Java Microservices with Apache Cassandra





# Your hosts



Cédrick Lunven



Developer Advocate



Eric Zietlow   
Developer Advocate



Jack Fryer



Community Manager



David Jones-Gilardi   
Developer Advocate



Aleksandr Volochnev   
Developer Advocate



# Spring Microservices with Apache Cassandra™

- 1 Housekeeping
- 2 Running the TodoApp
- 3 Connectivity to Cassandra
- 4 Create the Repository (dao)
- 5 Spring-Data & Spring-Data-Rest
- 6 Resources

# menti.com

47 86 87



Available on the iPhone  
**App Store**

GET IT ON  
**Google play**

# Spring Microservices with Apache Cassandra™

- 1 Housekeeping
- 2 Running the TodoApp
- 3 Connectivity to Cassandra
- 4 Create the Repository (dao)
- 5 Spring-Data & Spring-Data-Rest
- 6 Resources

# Disclaimer

Today, with limited time (2H), we won't introduce the basics of Apache Cassandra™.

This is a coding session. You would need a previous experience with Java language and the following tools :

## LOCAL Development

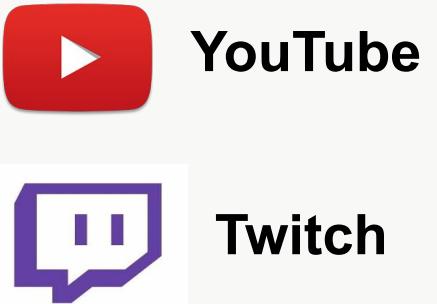
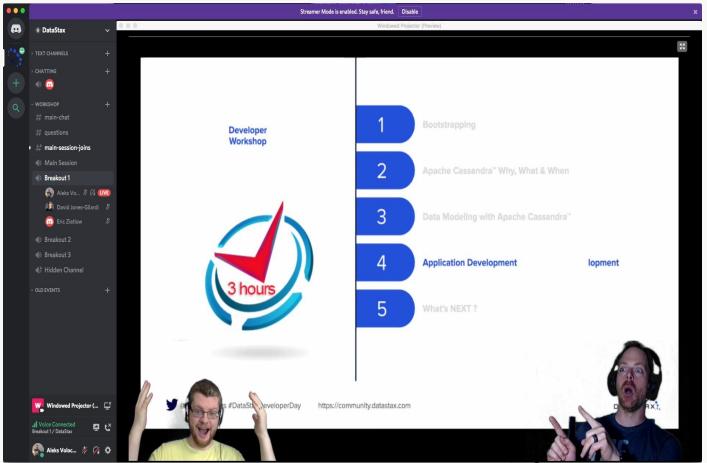
- ★ **JAVA (JDK8 +)**
- ★ **Maven**
- ★ **Git**
- ★ **an IDE**

OR

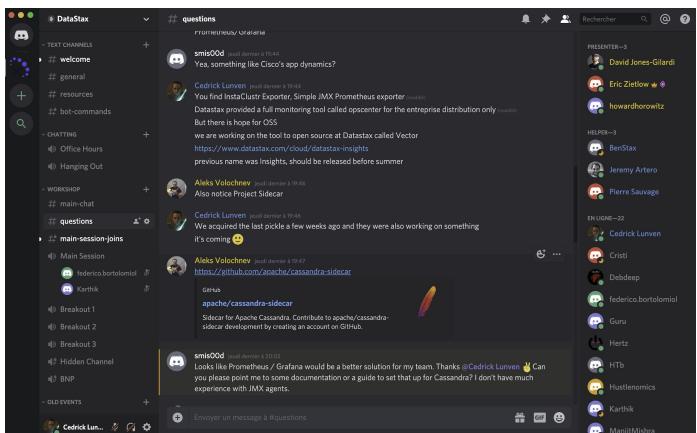
## GITPOD



## STREAMS

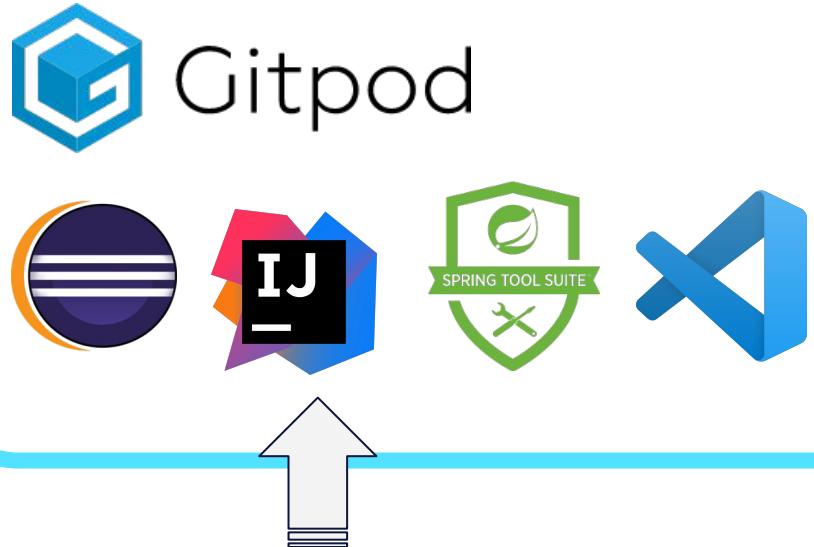


## QUESTIONS

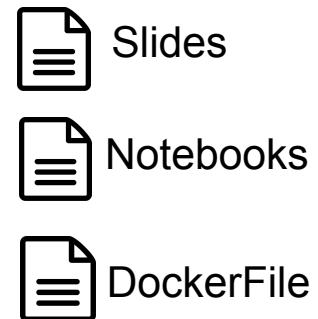


Discord

## RUNTIME



## MATERIALS



DATA  
STAX

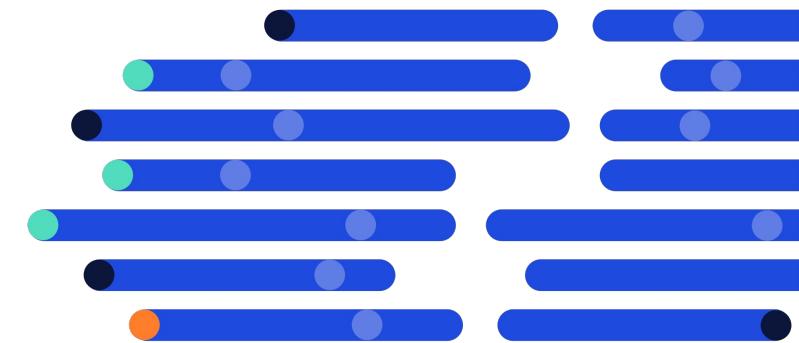
# Screen with youtube, discord, Astra & IDE

# Exercises

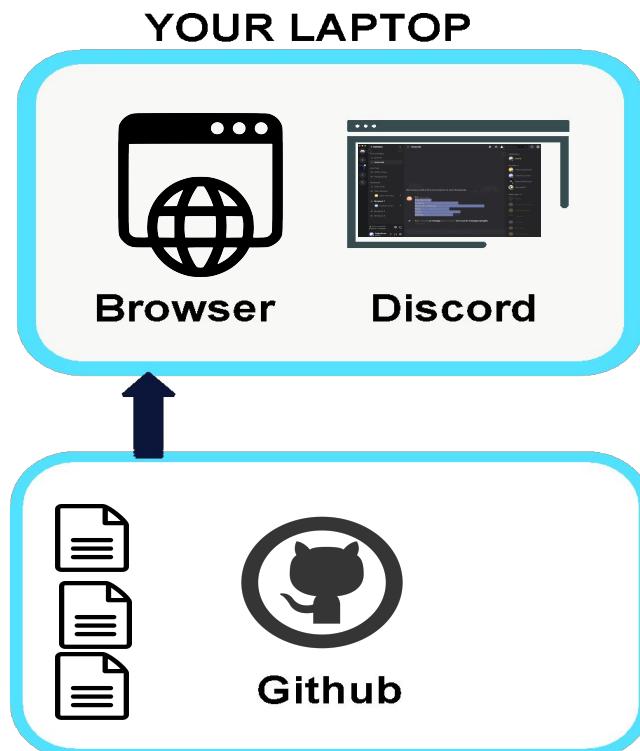
[https://github.com/DataStax-Academy/  
microservices-java-workshop-online](https://github.com/DataStax-Academy/microservices-java-workshop-online)

✨ Building a Microservices for Apache Cassandra™ ✨

Gitpod ready-to-code license Apache-2.0 chat 85 online



# Bootstrapping



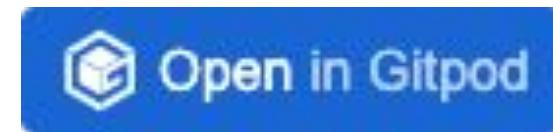
## 1. Clone or download repository material

<https://github.com/DataStax-Academy/microservices-java-workshop-online>

# Bootstrapping your env 2/2



## 3. Import the maven project in your IDE or open GitPod



## 4. Start Cassandra locally

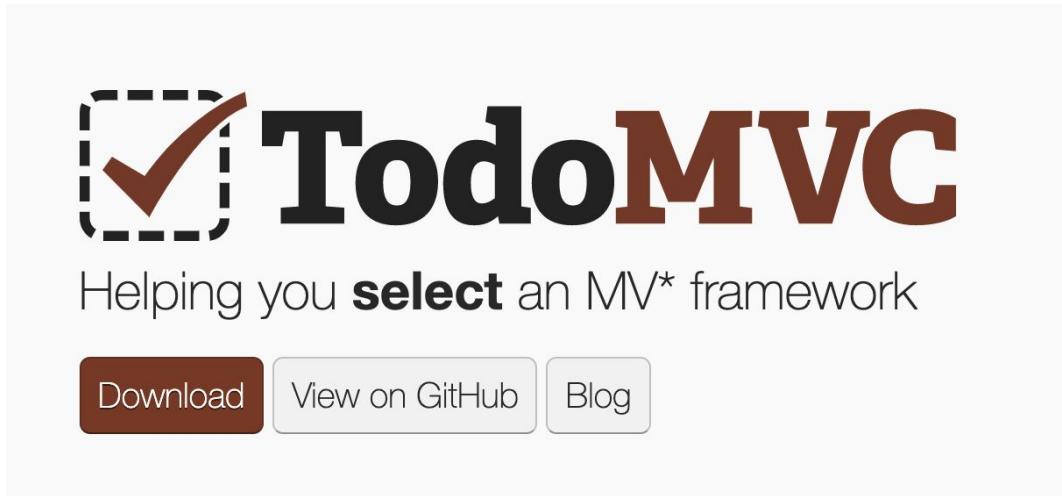
```
docker-compose up -d
```

# Spring Microservices with Apache Cassandra™

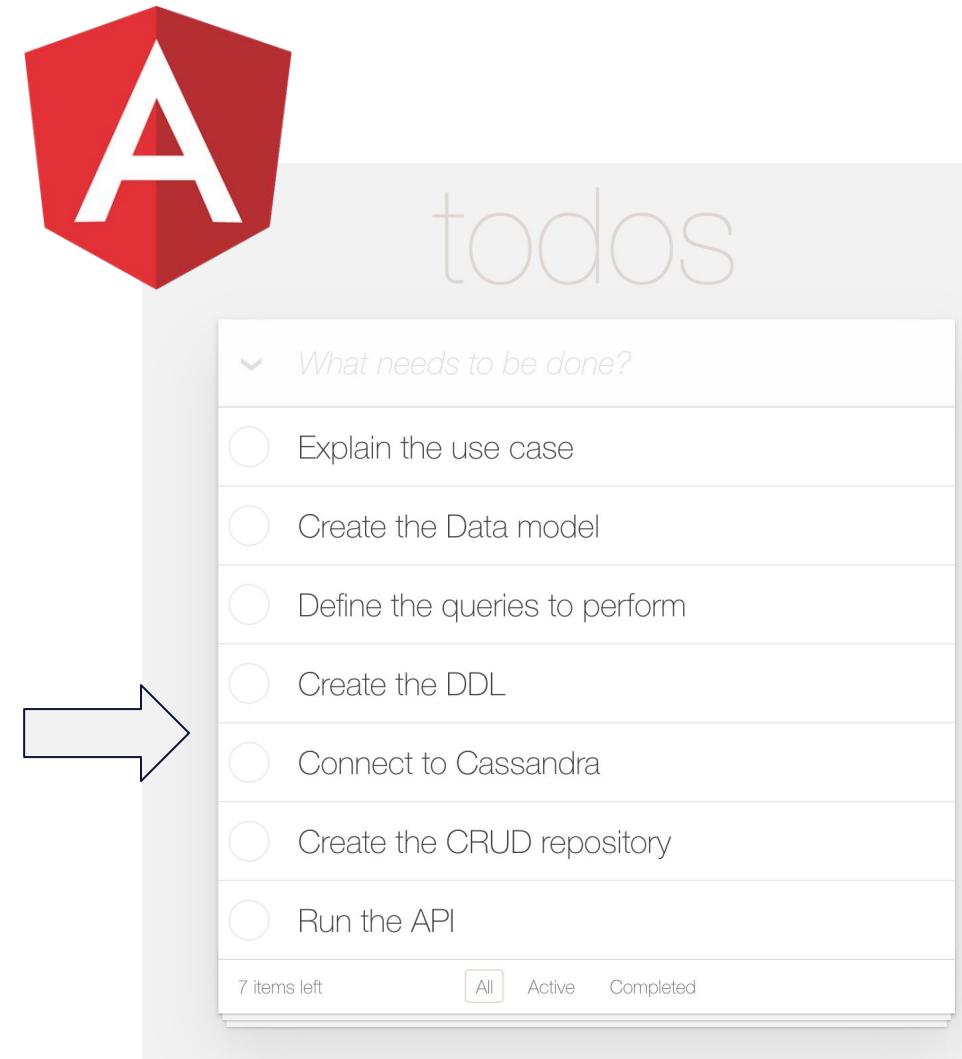
- 1 Housekeeping
- 2 Running the TodoApp
- 3 Connectivity to Cassandra
- 4 Create the Repository (dao)
- 5 Spring-Data & Spring-Data-Rest
- 6 Resources

# The Todo Application

<http://todomvc.com/>



<http://todomvc.com/examples/angularjs/>



# The Todo Application

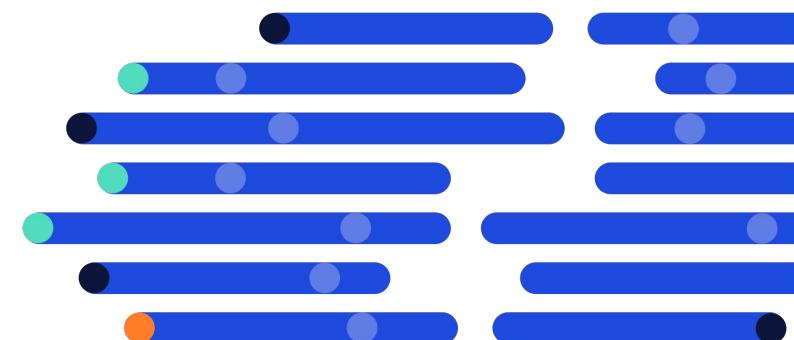


# Exercise 1a

---



- **Run the client on your machine**



# The Todo Application

<https://www.todobackend.com/>

# Todo-Backend

a shared example to showcase backend tech stacks

The Todo-Backend project defines a simple web API spec - for managing a todo list. Contributors [implement](#) that spec using various tech stacks. Those implementations are cataloged below. A spec runner verifies that each contribution implements the exact same API, by running an automated test suite which [defines](#) the API.

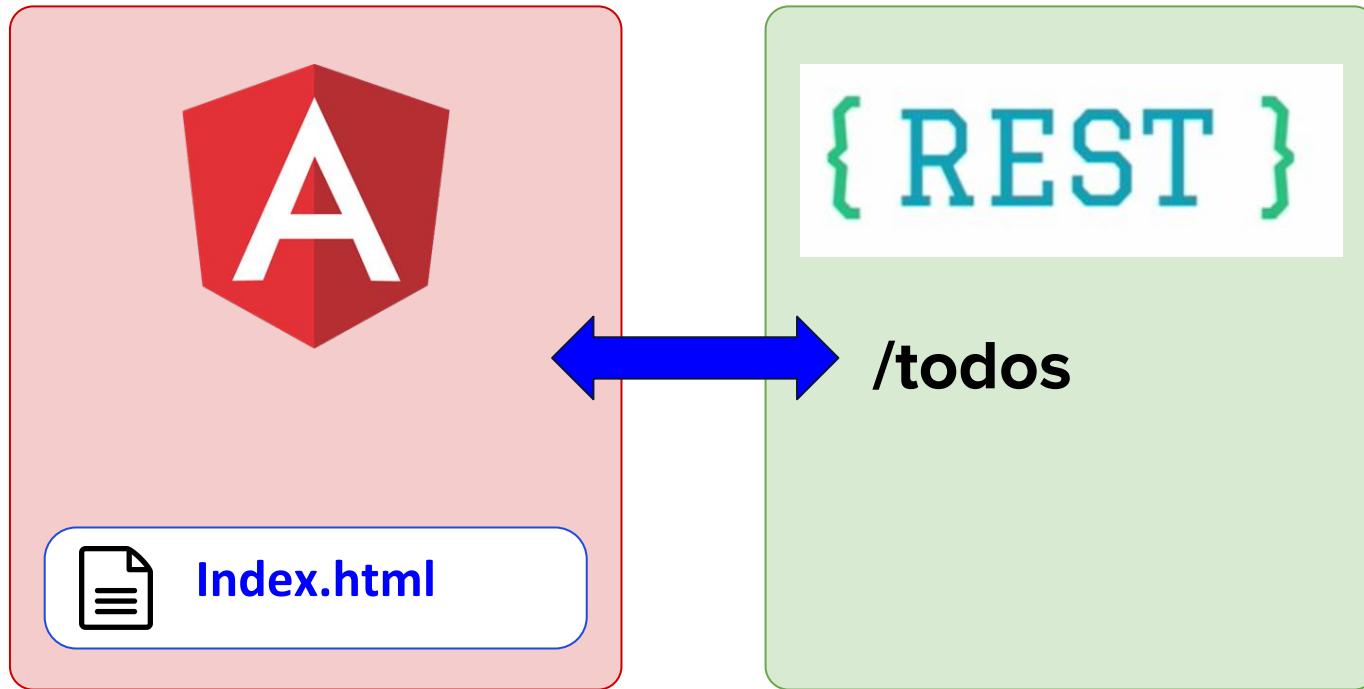
The Todo-Backend project was inspired by the [TodoMVC](#) project, and some code (specifically the todo client app) was borrowed directly from TodoMVC.

Created and curated by [Pete Hodgson](#).

featuring HTTP APIs built with:



# The Todo Application

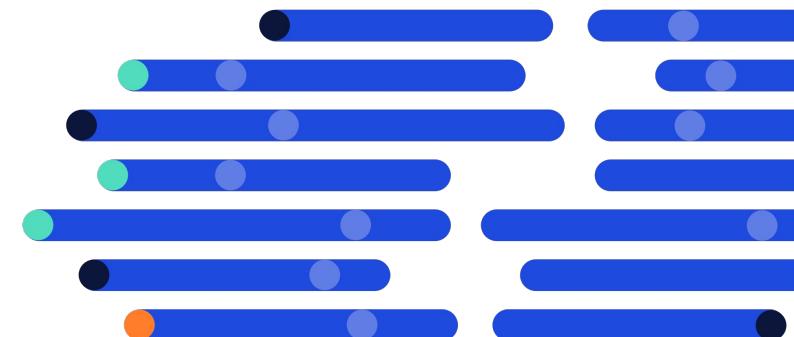


# Exercise 1b

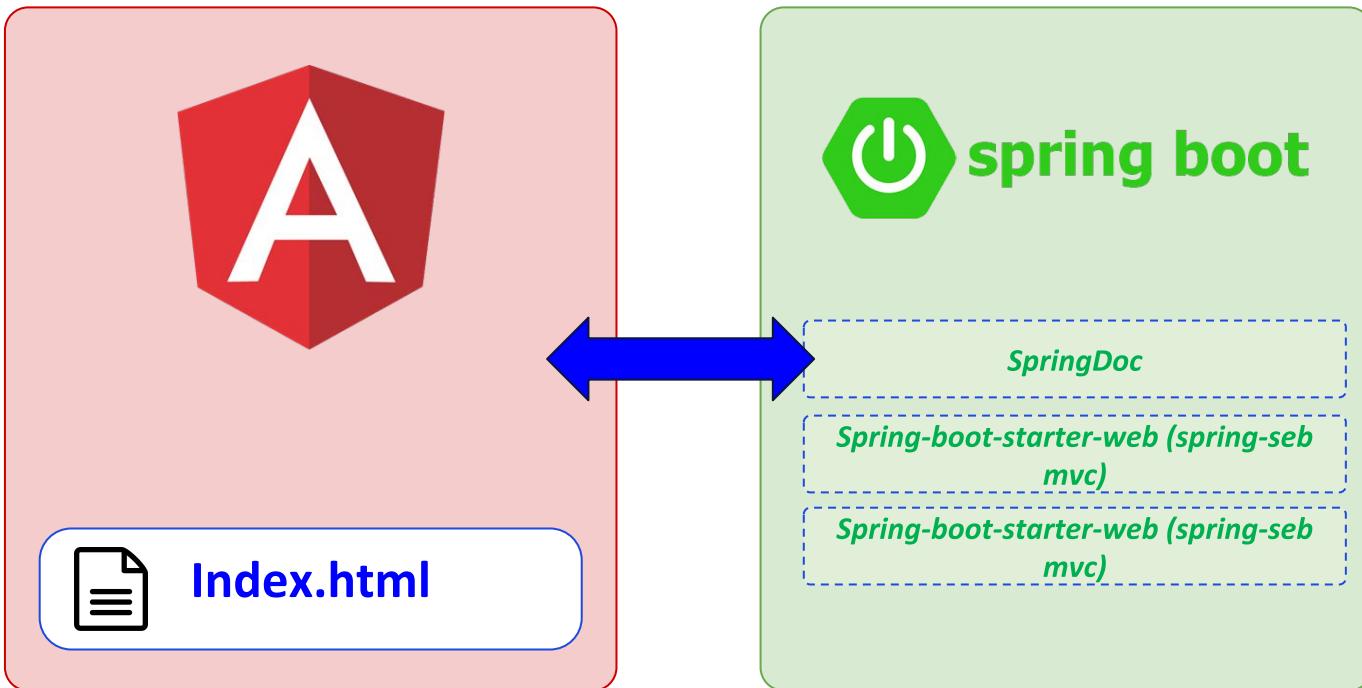
---



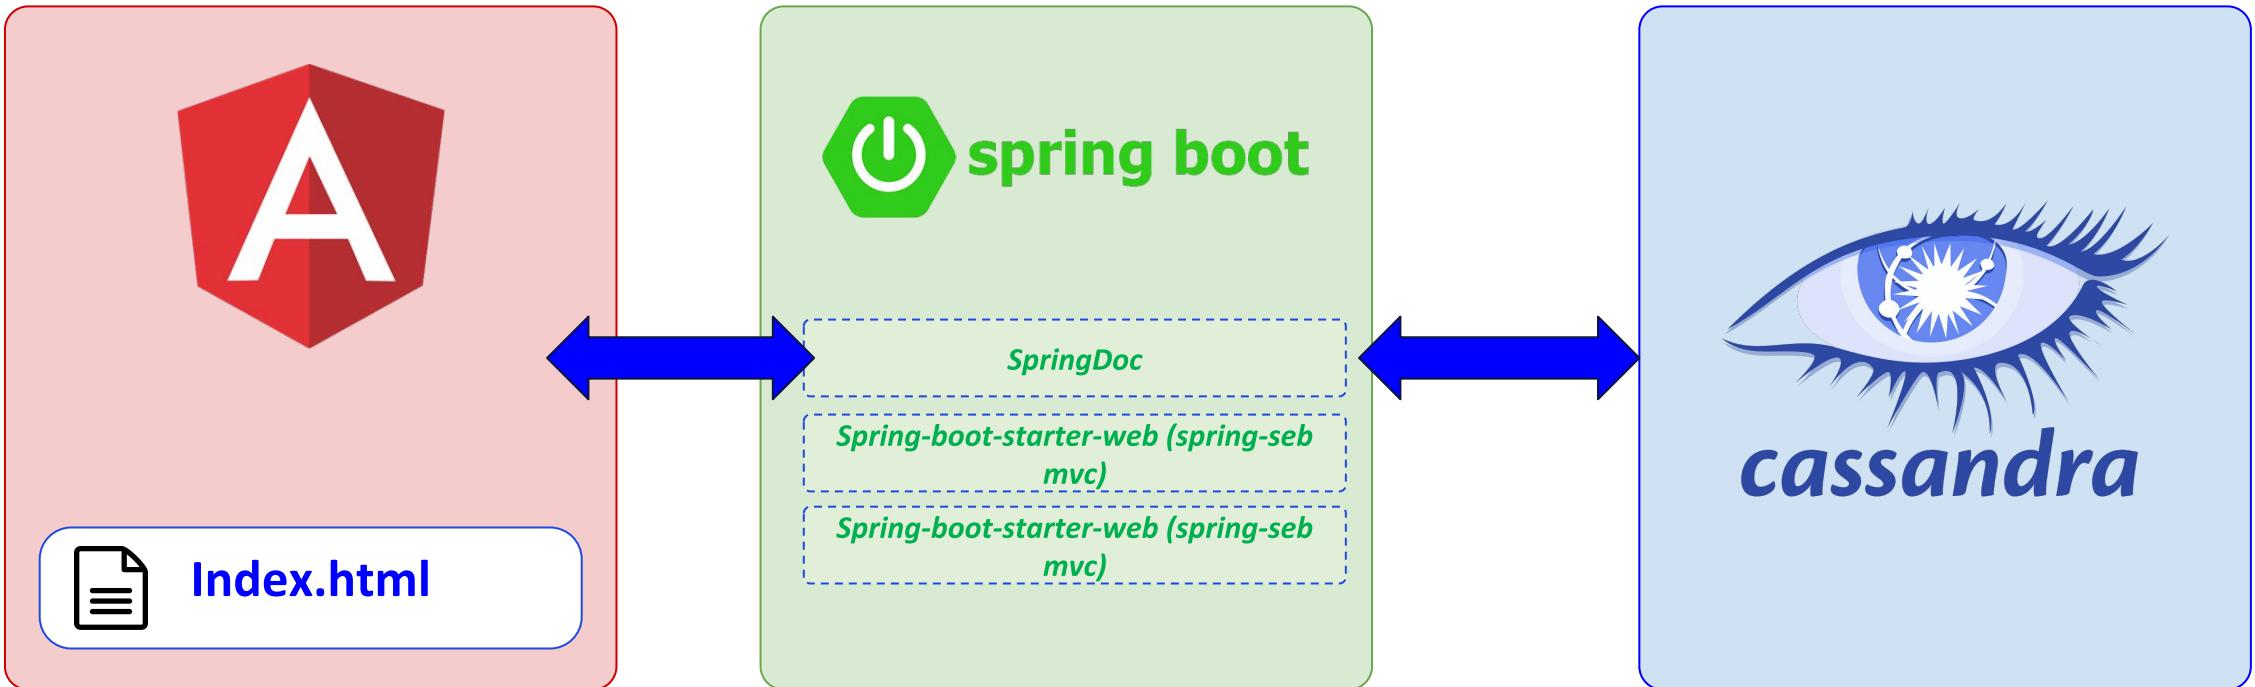
- **Start the backend**
- **Browse and test from generated documentation**
- **Test the API from the specification**



# The Todo Application



# The Todo Application





# cassandra ☆

Docker Official Images

Apache Cassandra is an open-source distributed storage system.

100M+

Container

Linux

PowerPC 64 LE

ARM

ARM 64

386

x86-64

Databases

Official Image

Linux - x86 ( latest )

Copy and paste to pull this image

`docker pull cassandra`



[View Available Tags](#)

## Running Cassandra in Docker

- Define a proper **network**
- **Env variables** can be defined to override keys in `cassandra.yaml`.
- Export ports **7000, 9042, ...**
- Define volumes to stores data
  - **/var/lib/cassandra**

```
$ docker run
--name some-cassandra -d \
-e CASSANDRA_BROADCAST_ADDRESS=10.42.42.42 \
-p 7000:7000,9042:9042
-v /my/own/datadir:/var/lib/cassandra \
cassandra:tag
```

```
docker-compose up -d
```

# Run Cassandra Locally

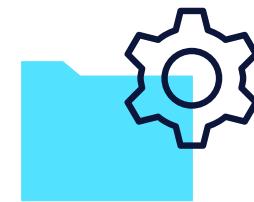
Define and run multi-container Docker applications through the use of a YAML file to configure your applications

```
version: '2'
services:

cassandra-seed:
  container_name: cassandra-seed-node
  image: cassandra: 3.11.6
  ports:
    - "9042:9042"      # Native transport
    - "7199:7199"      # JMX
    - "9160:9160"      # Thrift clients

cassandra-node:
  image: cassandra: 3.11.6
  command: /bin/bash -c "echo 'Waiting for seed node' && sleep 30 && /docker-entrypoint.sh cassandra -f"
  environment:
    - "CASSANDRA_SEEDS=cassandra-seed-node"
  depends_on:
    - "cassandra-seed"
```

# Introducing Astra



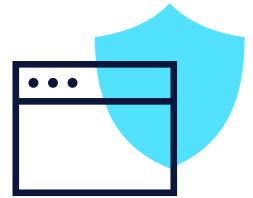
## Eliminate Operations

everything from provisioning to backups is fully automated



## Secure Your Data

with the most advanced security available for Cassandra



## Simplify App Development

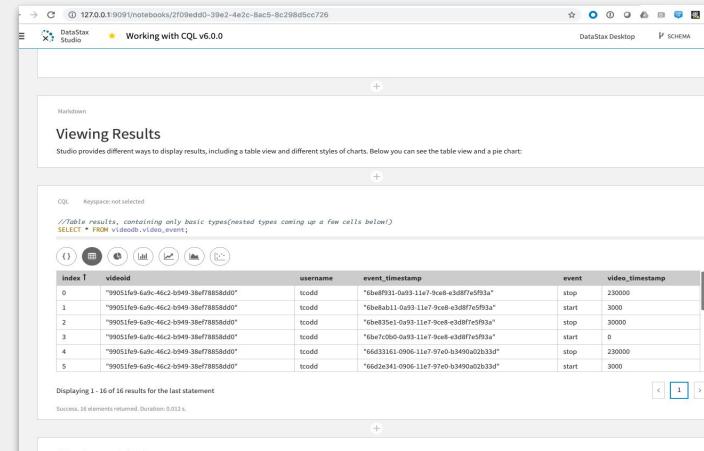
with auto-configured developer tools that deploy with a click

# Simplify Application Development

## Familiar Language

```
INSERT INTO mytable
(id, name, address) VALUES
(1, 'Bob Smith', '1 Main
Street')
SELECT * FROM mytable
WHERE id=1
UPDATE mytable SET
name='Tom Smith' WHERE
id=1
DELETE FROM mytable WHERE
id=1
```

## Easy Dev Tools



## Great Drivers



# Exercise 2



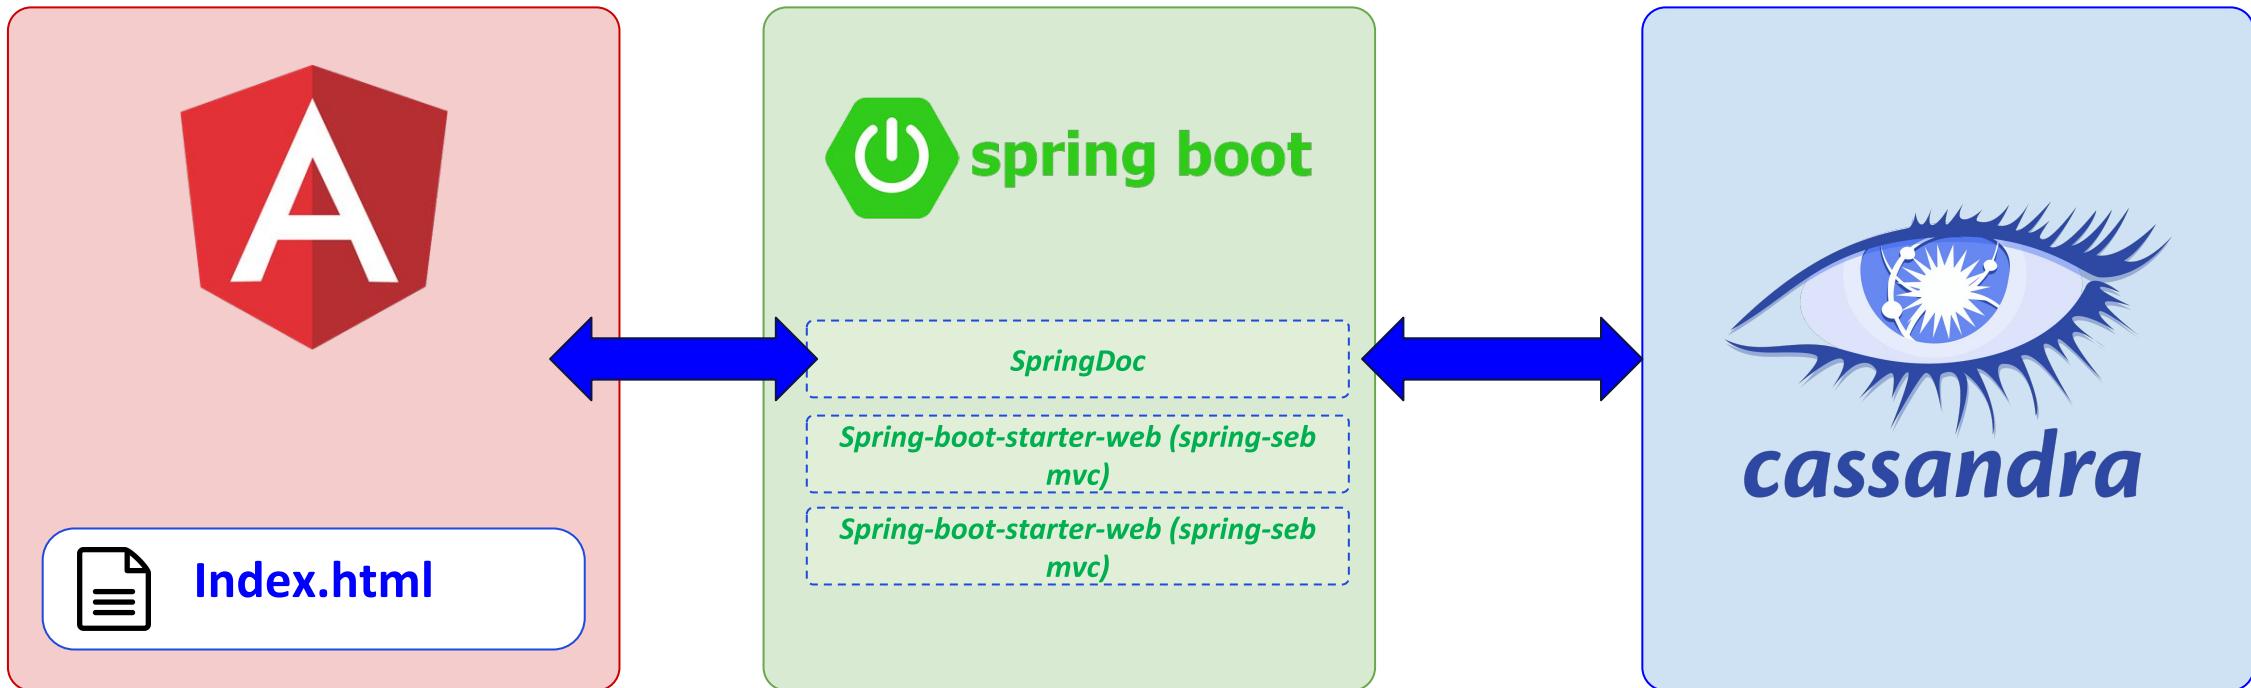
- Create your Astra Instance

The screenshot shows the DataStax Constellation web interface. At the top, there's a banner indicating "Database is initializing. You'll receive an email when the database is ready. While you're waiting, learn how to [get started with your database](#). Database creation started at Jan 22nd 2020 11:49am UTC. This can take as long as 30 minutes but is usually much quicker." Below this, the main dashboard shows a "screenshot" of a database named "okidoki" created by "Cedrick Lunven" on "January 22, 2020". The "Size and Location" section shows "1 Capacity Unit" and "500 GB Total Storage" across "us-east-1 (1 capacity unit)". The "Cost" section shows "Spent this month: TBD" and "Estimated Cost" options for "Per Hour" and "Per Month". The "Connection Details" section notes that connection details are only available for active databases. The top right corner shows the user "cedrick.lunven@datastax.com".

# Spring Microservices with Apache Cassandra™

- 1 Housekeeping
- 2 Running the TodoApp
- 3 Connectivity to Cassandra
- 4 Create the Repository (dao)
- 5 Spring-Data & Spring-Data-Rest
- 6 Resources

# DataStax Drivers



# DataStax Drivers Features

*One of set drivers to connect them all - january  
2020*



## Connectivity

- ★ Token & Datacenter Aware
- ★ Load Balancing Policies
- ★ Retry Policies
- ★ Reconnection Policies
- ★ Connection Pooling
- ★ Health Checks
- ★ Authentication | Authorization
- ★ SSL

## Query

- ★ CQL Support
- ★ Schema Management
- ★ Sync/Async/Reactive API
- ★ Query Builder
- ★ Compression
- ★ Paging

## Parsing Results

- ★ Lazy Load
- ★ Object Mapper
- ★ Spring Support
- ★ Paging

# File-based Configuration

- Based on Typesafe Config
- Attributes are grouped into basic and advanced categories
- A reference file (*reference.conf*) provide default values embedded in the jar file. Can be override with key in application.conf.
- Driver searches `application.conf` in the classpath



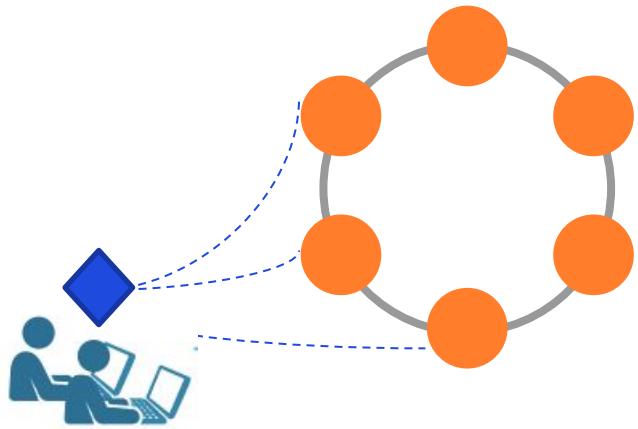
`application.conf`

```
datastax-java-driver {  
    basic {  
        request.timeout      = 5 seconds  
        request.consistency = LOCAL_QUORUM  
    }  
}
```



# Load-Balancing

- Used to create query plans for each statement executed
- Default policy is token-aware, round robin
- Requests are routed to nodes in the “local” data center only



## application.conf

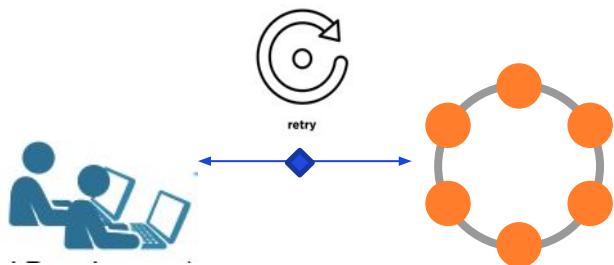
```
datastax-java-driver {  
    basic {  
        load-balancing-policy {  
  
            # The class of the policy.  
            class = DefaultLoadBalancingPolicy  
  
            # The datacenter that is considered "local"  
            # The default policy will only include nodes from  
            # this datacenter in its query plans.  
            local-datacenter = datacenter1  
  
            # A custom filter to include/exclude nodes  
            // filter.class=  
        }  
    }  
}
```



# Retry Policy

Determines when queries are retried on failure

- **DefaultRetryPolicy**
  - Default
  - Retries once onReadTimeout or onWriteTimeout
  - Enough replicas for your consistency level must be online
  - Only retries idempotent mutations



application.conf

```
datastax-java-driver {  
    # The policy that controls if the driver retries  
    # requests that have failed on one node.  
    advanced.retry-policy {  
        # The class of the policy  
        class = DefaultRetryPolicy  
    }  
}
```



@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>

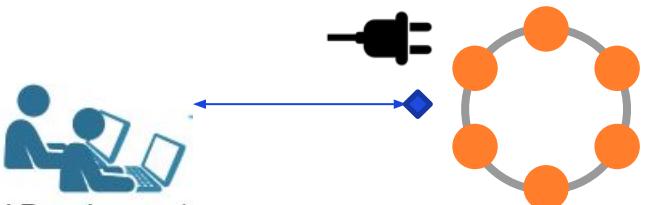


# Reconnection Policy

Reconnects driver to a downed node

Two options:

- **ConstantReconnectionPolicy**
  - Check every N milliseconds
- **ExponentialReconnectionPolicy**
  - Increases every interval
  - Caps out at a max



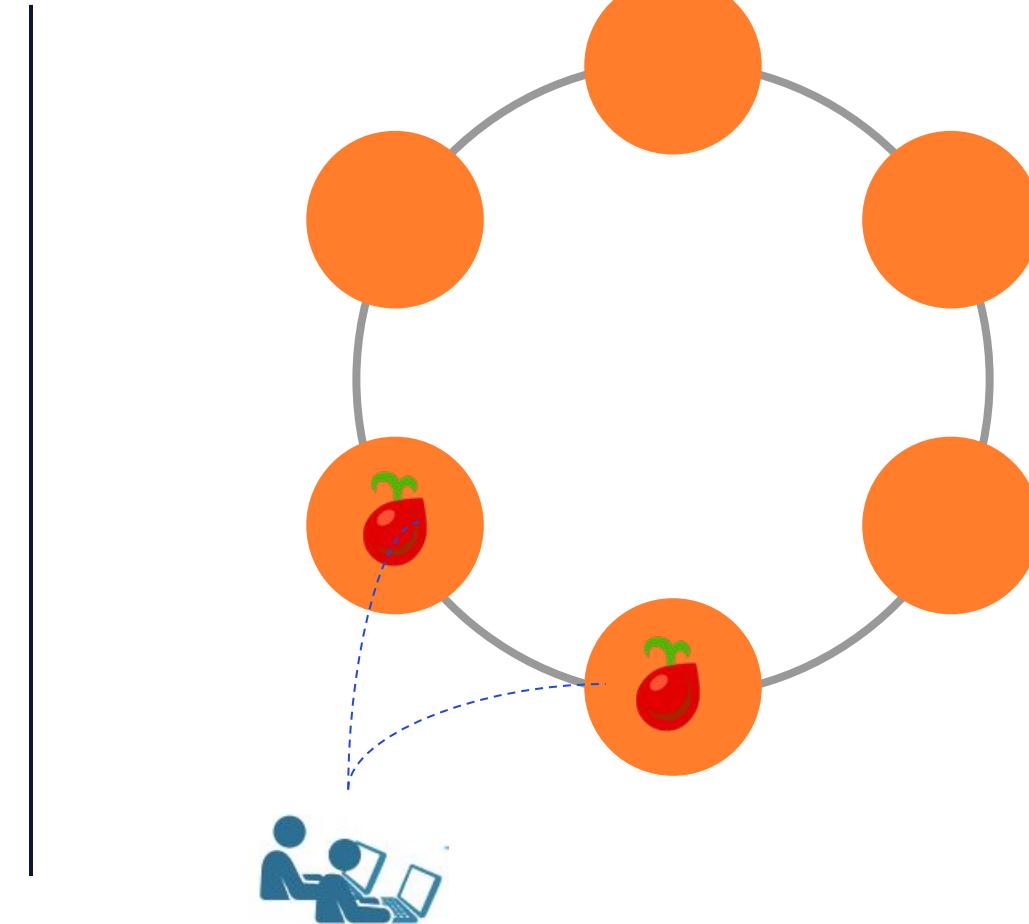
application.conf

```
datastax-java-driver {  
  
    # Whether to schedule reconnection attempts  
    # if all contact points are unreachable at init  
    advanced.reconnect-on-init = false  
  
    advanced.reconnection-policy {  
  
        # The class of the policy  
        class = ExponentialReconnectionPolicy  
  
        # Parameters  
        base-delay = 1 second  
        max-delay = 60 seconds  
    }  
}
```



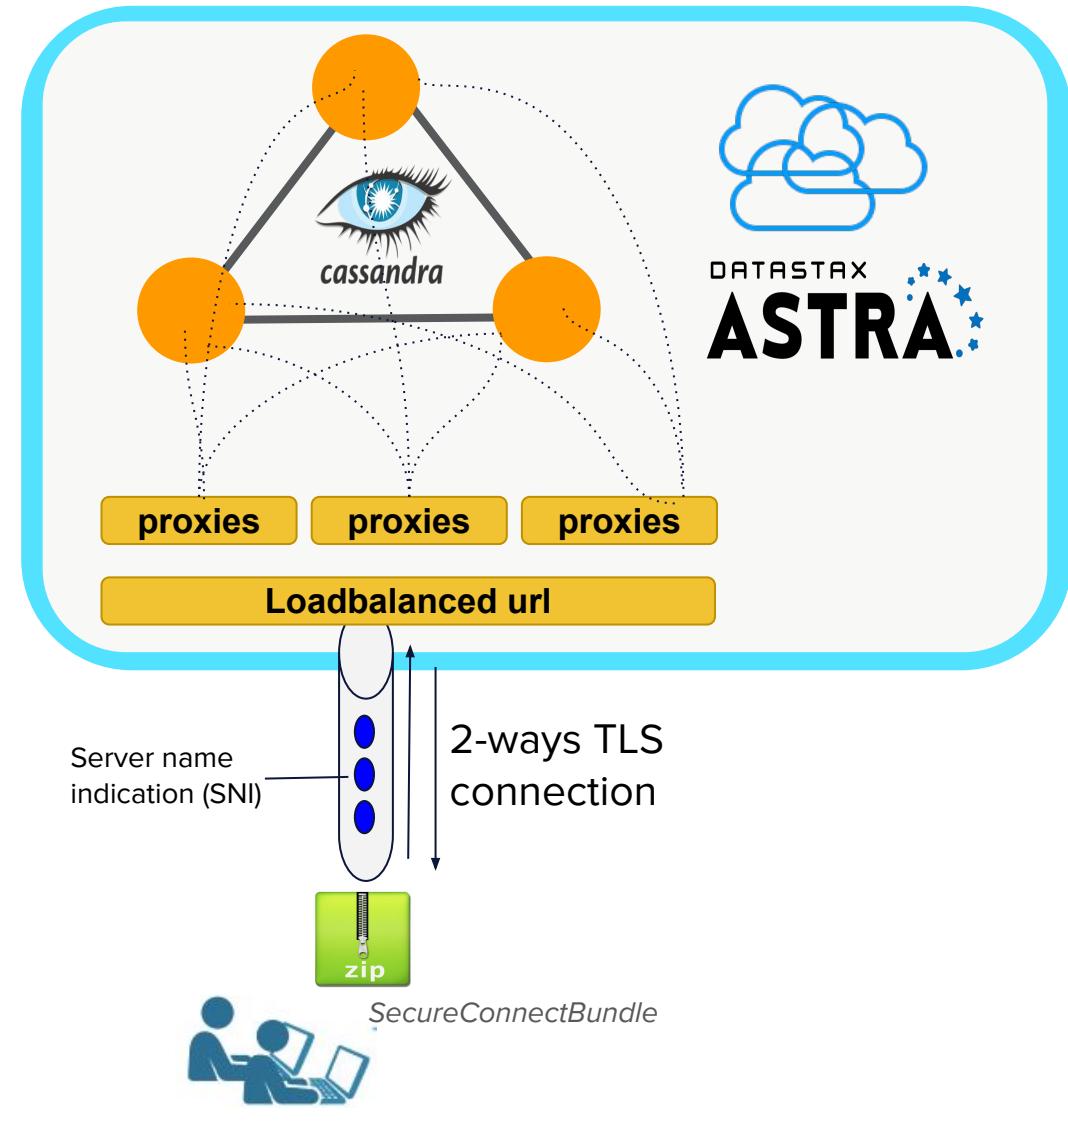
# Contact Points

- Only one necessary
- Unless that node is down
- More are good



# Contact Points with ASTRA

- Download the SecureConnectBundle from Astra UI.
- Provide **username, password and keyspace name in the BUILDER.**



# Dependencies



# Open a Session

```
// Delegate all configuration to file or default
CqlSession cqlSession = CqlSession.builder().build();

// Explicit Settings
CqlSession cqlSession = CqlSession.builder()
    // .addContactPoint(new InetSocketAddress("127.0.0.1", 9042))
    .withCloudSecureConnectBundle(Paths.get("/tmp/apollo.zip"))
    .withContactPoint
    .withKeyspace("killrvideo")
    .withAuthCredentials("KVUser", "KVPassword")
    .build();
```

# Important to know about CqlSession

- **CqlSession** is a stateful object handling communications with each node
- **CqlSession** should be unique in the Application (*Singleton*)
- **CqlSession** should be closed at application shutdown (*shutdown hook*) in order to free opened TCP sockets (*stateful*)

```
@PreDestroy  
public void cleanup() {  
    if (null != cqlSession) {  
        cqlSession.close();  
    }  
}
```

# Exercise 3



- **Download secure bundle**
- **Make the Tests Green**
- **Connect, list keyspaces**

A screenshot of the DataStax Constellation web interface. At the top, it says "Your Organization". Below that, there's a green banner stating "Database is initializing. You'll receive an email when the database is ready. While you're waiting, learn how to get started with your database." The main content area shows a table with two columns: "screenshot" and "Size and Location". The "screenshot" column contains information about the keyspace: "Keyspace Name: okidoki", "Organization: Your Organization", "Owner: Cedrick Lunven", and "Created: January 22, 2020". The "Size and Location" column shows "1 Capacity Unit", "500 GB Total Storage", "Locations: us-east-1 (1 capacity unit)", "Compute Size: Startup", and "Replication Factor: 3". Below this, there are two more sections: "Cost" and "Connection Details". The "Cost" section shows "Spent this month: TBD" and "Estimated Cost" with options for "Per Hour" and "Per Month". The "Connection Details" section notes that details are only available for active databases.

screenshot	Size and Location
Keyspace Name: okidoki Organization: Your Organization  Owner: Cedrick Lunven Created: January 22, 2020	1 Capacity Unit 500 GB Total Storage  Locations: us-east-1 (1 capacity unit) Compute Size: Startup Replication Factor: 3

Cost	Connection Details
Spent this month: TBD  Estimated Cost <input checked="" type="radio"/> Per Hour <input type="radio"/> Per Month  when running TBD/hour when parked TBD/hour	Connection details are only available for active databases that you own or have connection permissions for.

# Spring Microservices with Apache Cassandra™

- 1 Housekeeping
- 2 Running the TodoApp
- 3 Connectivity to Cassandra
- 4 Create the Repository (dao)
- 5 Spring-Data & Spring-Data-Rest
- 6 Resources

# Build Proper Data Model

We know the requests from the specs

# How to execute queries ?

- First job of **CqlSession** is to execute queries using, well, execute method.

```
cqlSession.execute("SELECT * FROM killrvideo.users");
```

Statement

# SimpleStatement

```
Statement statement = ...  
  
// (1) Explicit SimpleStatement Definition  
SimpleStatement.newInstance("select * from t1 where c1 = 5");  
  
// (2) Externalize Parameters (no name)  
SimpleStatement.builder("select * from t1 where c1 = ?")  
    .addPositionalValue(5);  
  
// (3) Externalize Parameters (name)  
SimpleStatement.builder("select * from t1 where c1 = :myVal")  
    .addNamedValue("myVal", 5);  
  
cqlSession.execute(statement);
```

# Prepared and Bound Statements

- Compiled once on each node automatically as needed
- Prepare each statement only once per application
- Use one of the many bind variations to create a BoundStatement

```
PreparedStatement ps = cqlSession.prepare("SELECT * from t1 where c1 = ?");  
  
BoundStatement bound = ps.bind(5);  
  
cqlSession.execute(bound);
```

# ResultSet

- **ResultSet** is the object returned for executing query. It contains **ROWS** (data) and **EXECUTION INFO**.
- **ResultSet** is **iterable** and as such you can navigate from row to row.
- Results are **always paged** for you (avoiding memory and response time issues)

```
ResultSet rs = cqlSession.execute(myStatement);

// Plumbery
ExecutionInfo info = rs.getExecutionInfo();
int executionTime = info.getQueryTrace().getDurationMicros();

// Data: NOT ALL DATA RETRIEVED IMMEDIATELY (only when needed .next())
Iterator<Row> iterRow = rs.iterator();
int itemsFirstCall = rs.getAvailableWithoutFetching();
```

# Parsing ResultSet

```
// We know there is a single row (eg: count)
Row singleRow = resultSet.one();

// We know there are not so many results we can get all (fetch all pages)
List<Row> allRows = resultSet.all();

// Browse iterable
for(Row myRow : resultSet.iterator()) {
    // .. Parsing rows
}

// Use Lambda
rs.forEach(row -> { row.getColumnDefinitions(); });

// Use for LWT
boolean isQueryExecuted = rs.wasApplied();
```

# Parsing Rows

```
// Sample row
Row row = resultSet.one();

// Check null before read
Boolean isUserNameNull = row.isNull("userName");

// Reading Values from row
String userName1 = row.get("username", String.class);
String userName2 = row.getString("username");
String userName3 = row.getString(CqlIdentifier.fromCql("username"));

// Tons of types available
row.getUuid("userid");
row.getBoolean("register");
row.getCqlDuration("elapsed");
...
```

# Paging

- ResultSet contains up to “[pageSize](#)” items. When browsing records you may hit this number that will trigger fetching next “pageSize” items.
- To fetch anything else than first page you must provide a [PagingState](#).

```
// Enforce few items per page (often = UI requirements)
myStatement = myStatement.setPageSize(10);
ResultSet page1 = cqlSession.execute(myStatement);

// Paging State
ByteBuffer pagingState = page1.getExecutionInfo().getPagingState();
myStatement = myStatement.setPageState(pagingState);

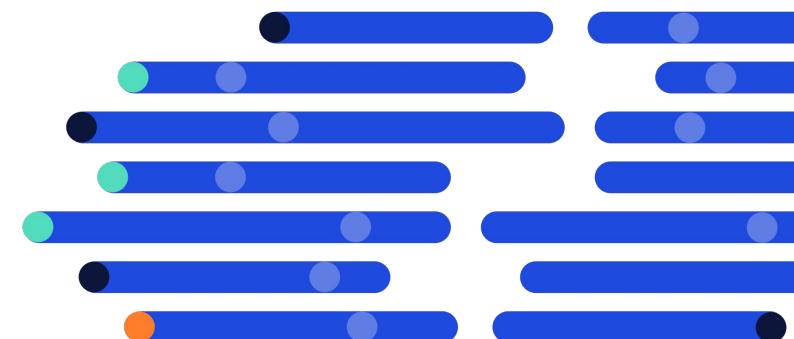
// Very same statement with pagingState provided
ResultSet page2 = cqlSession.execute(myStatement);
```

# Exercise 4a

---



- **CRUD Repository Ad Hoc**



# QueryBuilder

- Fluent API for building CQL string queries programmatically
- Contains methods to build SELECT, UPDATE, INSERT and DELETE statements
- Generates a Statement as per the earlier techniques

## OSS Driver (current version 4.2.0)

```
<dependency>
  <groupId>com.datastax.oss</groupId>
  <artifactId>java-driver-query-builder</artifactId>
</dependency>
```



@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# QueryBuilder

```
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.bindMarker;
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.deleteFrom;
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.selectFrom;
import static com.datastax.oss.driver.api.querybuilder.relation.Relation.column;

// Simple SELECT using QueryBuilder
Statement stmtSelect = selectFrom("killrvideo", "videos_by_users")
    .column("userid").column("commentid")
    .function("toTimestamp", Selector.column("commentid")).as("comment_timestamp")
    .where(column("userid").isEqualTo(bindMarker("userid"))))
    .build()

// Simple DELETE using QueryBuilder
Statement stmtDelete = deleteFrom("killrvideo", "videos_by_users")
    .where(column("userid").isEqualTo(bindMarker("userid"))))
    .build()
```



# QueryBuilder

- Can also use **QueryBuilder** to create **PreparedStatements** and later execute at runtime
- Note use of **bindMarker()** to designate parameters that will be provided later

```
// Prepared QueryBuilder statements as any statement
PreparedStatement psStmt = cqlSession.prepare(
    deleteFrom("killrvideo", "videos_by_users")
        .where(column("userid").isEqualTo(bindMarker("userid"))))
        .build()));

// Binding
BoundStatement bsStmt = psStmt.bind("e7a8ac9f-c12d-415c-a526-4137815df573");

// Execute
cqlSession.execute(bsStmt);
```



# Batch Example

```
// Sample statements (insert same data in multiple tables)
Statement stmt1 = SimpleStatement
    .builder("INSERT INTO users_by_group(groupid,userid) values(?,?)")
    .addPositionalValue(groupname, username);
Statement stmt2 = SimpleStatement
    .builder("INSERT INTO groups_by_user(userid,groupid) values(?,?)")
    .addPositionalValue(username, groupname);

// Group as a Batch
BatchStatement batchStmt = BatchStatement
    .builder(DefaultBatchType.LOGGED)
    .addStatement(stmt1).addStatement(stmt2).build();

// Execute
cqlSession.execute(batchStmt);
```

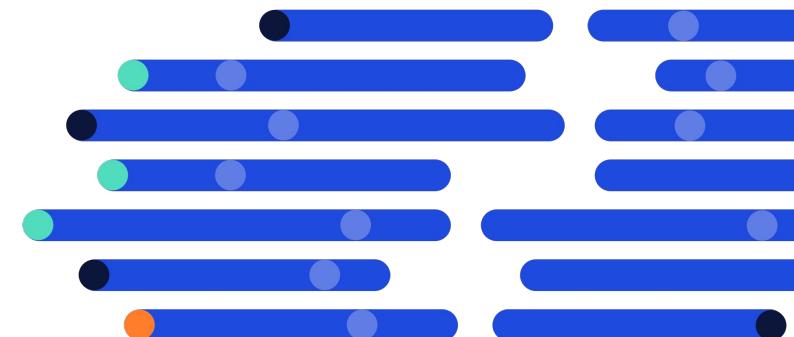


# Exercise 4b

---



- **CRUD Repository QueryBuilder**



# Object Mapper

- WHAT ?
  - Abstracts details of mapping Java attributes to/from CQL types and UDTs
  - Packaged separately from the driver – pom.xml update required
  - This slide shows the runtime dependency – will show compile-time shortly
- HOW ?
  - Some annotation processors will GENERATE Mapper, Dao, and Entity implementations for you
  - At each update in the files, the IDE (eclipse, intelliJ) will use annotation processor
  - Compiler plugin must be updated to define the annotation processor



# Object Mapper

## OSS Driver

```
<dependency>
  <groupId>com.datastax.oss</groupId>
  <artifactId>java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.oss</groupId>
      <artifactId>java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
  </configuration>
</plugin>
```

## DSE Driver

```
<dependency>
  <groupId>com.datastax.dse</groupId>
  <artifactId>dse-java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.dse</groupId>
      <artifactId>dse-java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
  </configuration>
</plugin>
```



# Annotate Entities

```
@Entity  
@CqlName("user_v")  
public class UserVideo {  
  
    @PartitionKey  
    @CqlName("userid")  
    private UUID userid;  
  
    @ClusteringColumn(1)  
    @CqlName("added")  
    private UUID videoid;  
  
}
```

TABLE NAME, KEYSPACE

PARTITION KEY COLUMNS

CLUSTERING COLUMNS



# Annotate DAO Interface (1/2)

```
@Dao
public interface VideoDao {

    @Select
    Optional<UserVideo> findUserById(UUID userid);

    @Query("SELECT * FROM ${tableId}")
    PagingIterable<UserVideo> findAll();

    @Select(customWhereClause = "videoid = : vid")
    PagingIterable<UserVideo>
    findUserByVideoId(@CqlName("videoid") UUID vid);
}
```



## Annotate DAO Interface (2/2)

```
// Save a bean
@Insert
void save(UserVideo userVideo);

// Userid id is PK
@Delete
void delete(UUID userid);

// Custom implementations
@QueryProvider(
    providerClass = MySampleQueryProvider.class,
    entityHelpers = { UserVideo.class })
String doSomething(String abc);
```



# Annotate Mapper Interface

```
@Mapper  
public interface MyApplicationMapper {  
  
    @DaoFactory  
    VideoDao videoDao(@DaoKeyspace CqlIdentifier keyspace);  
  
}
```



# Sample QueryProvider

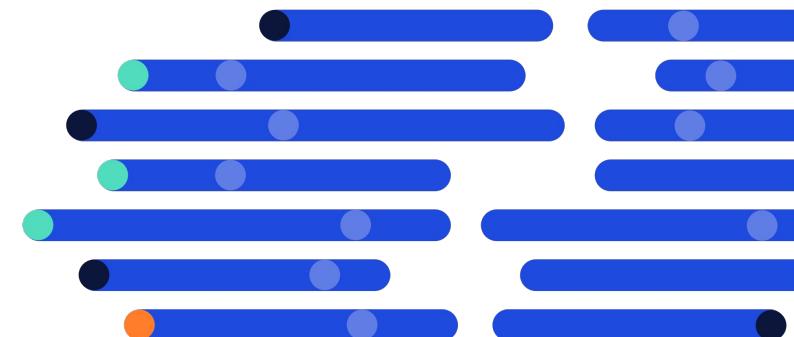
```
public class MySampleQueryProvider {  
  
    // Constructor, getting session  
    public MySampleQueryProvider(  
        MapperContext context,  
        EntityHelper<UserVideo> helperUser) {}  
  
    // Custom implementation method  
    public String doSomething(String abc) {  
    }  
}
```



# Exercise 4c



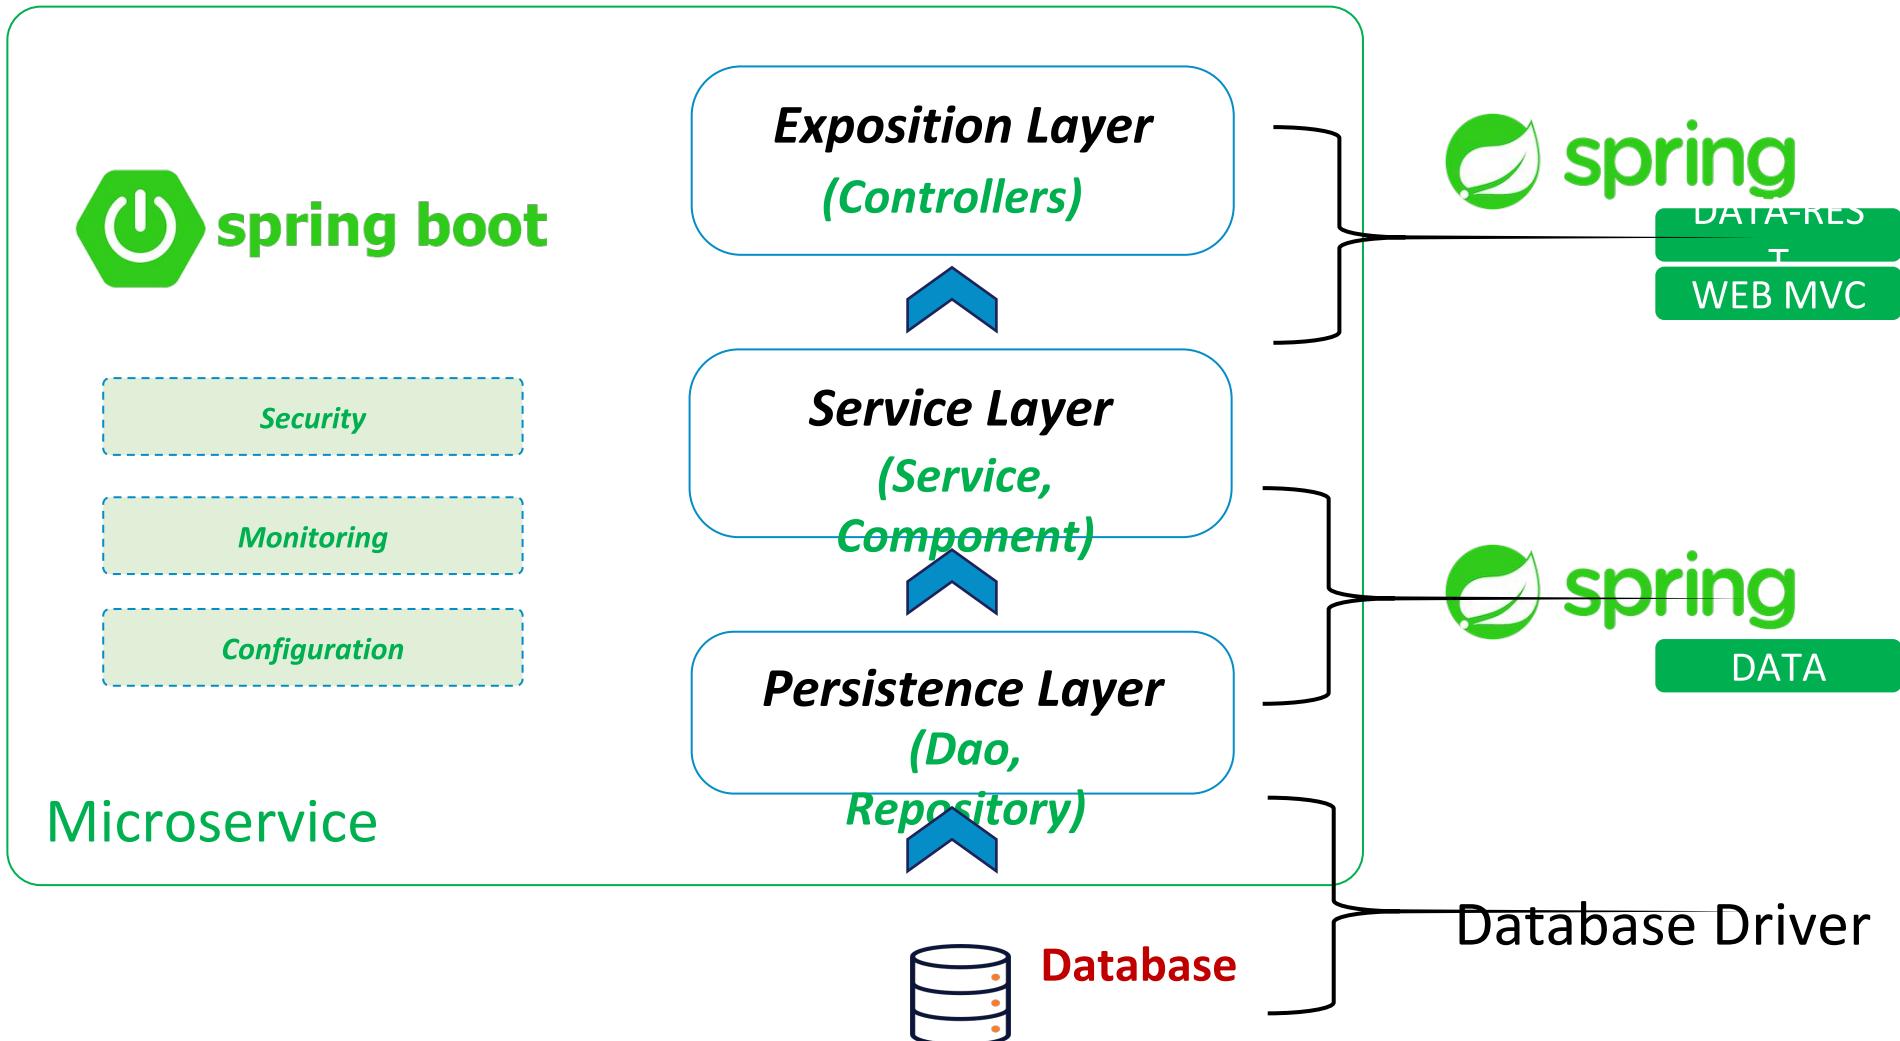
- **CRUD Repository Object Mapper**



# Spring Microservices with Apache Cassandra™

- 1 Housekeeping
- 2 Running the TodoApp
- 3 Connectivity to Cassandra
- 4 Create the Repository (dao)
- 5 Spring-Data & Spring-Data-Rest
- 6 Resources

« Vanilla » So easy....



## Entity => Spring-Data => Spring-Data-REST

```
@Entity  
public class Customer {  
  
    @Id  
    @GeneratedValue(strategy=GenerationType.AUTO)  
    private Long id;  
  
    private String firstName;  
  
    private String lastName;  
  
    private Customer() {}  
  
    private Customer(  
        String firstName, String lastName) {  
        this.firstName = firstName;  
        this.lastName = lastName;  
    }  
}
```

```
@RepositoryRestResource(  
    collectionResourceRel = "customers",  
    path = "customers")  
  
public interface CustomerRepository extends  
CrudRepository<Customer, Long> {  
  
    List<Customer> findByLastName(String lastName);  
  
    Customer findById(long id);  
}
```

## Configuration – 1# Convention

```
# -----
# Full Convention
# -----
spring:
  data:
    cassandra:
      contact-points: localhost
      port: 9042
      local-datacenter: dc1
      keyspace-name: betterbotz
      schema-action: create-if-not-exists
```

## Configuration – #2 Java Config

```
@Configuration
public class SpringDataCassandraJavaConfig
    extends AbstractCassandraConfiguration
    implements CqlSessionBuilderCustomizer {

    @Override
    protected String getKeyspaceName() {
        return keyspaceName;
    }

    @Override
    protected String getLocalDataCenter() {
        return localDataCenter;
    }
}
```

## Configuration - #3 Custom Definition

```
@Configuration
public class DseConfiguration {

    /** Internal logger. */
    private static final Logger LOGGER = LoggerFactory.getLogger(DseConfiguration.class);

    @Value("#{${dse.contactPoints}.split(',')}")
    public List < String > contactPoints;

    @Bean
    public DseSession dseSession() {
        long top = System.currentTimeMillis();
        LOGGER.info("Initializing connection to DSE Cluster");

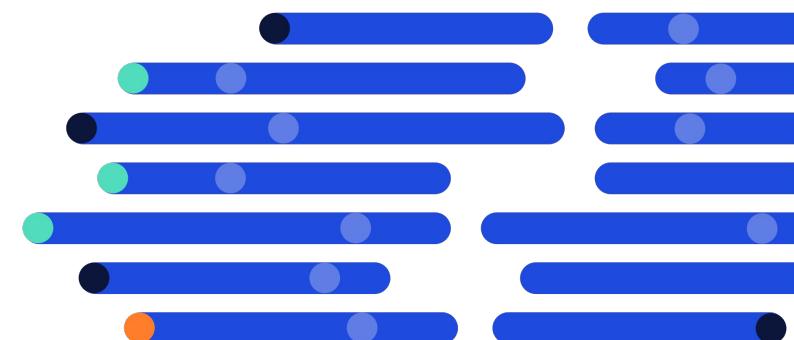
        Builder clusterConfig = new Builder();
        LOGGER.info(" + Contact Points : {}" , contactPoints);
        contactPoints.stream().forEach(clusterConfig::addContactPoint);
        LOGGER.info(" + Listening Port : {}", port);
        clusterConfig.withPort(port);
    }
}
```

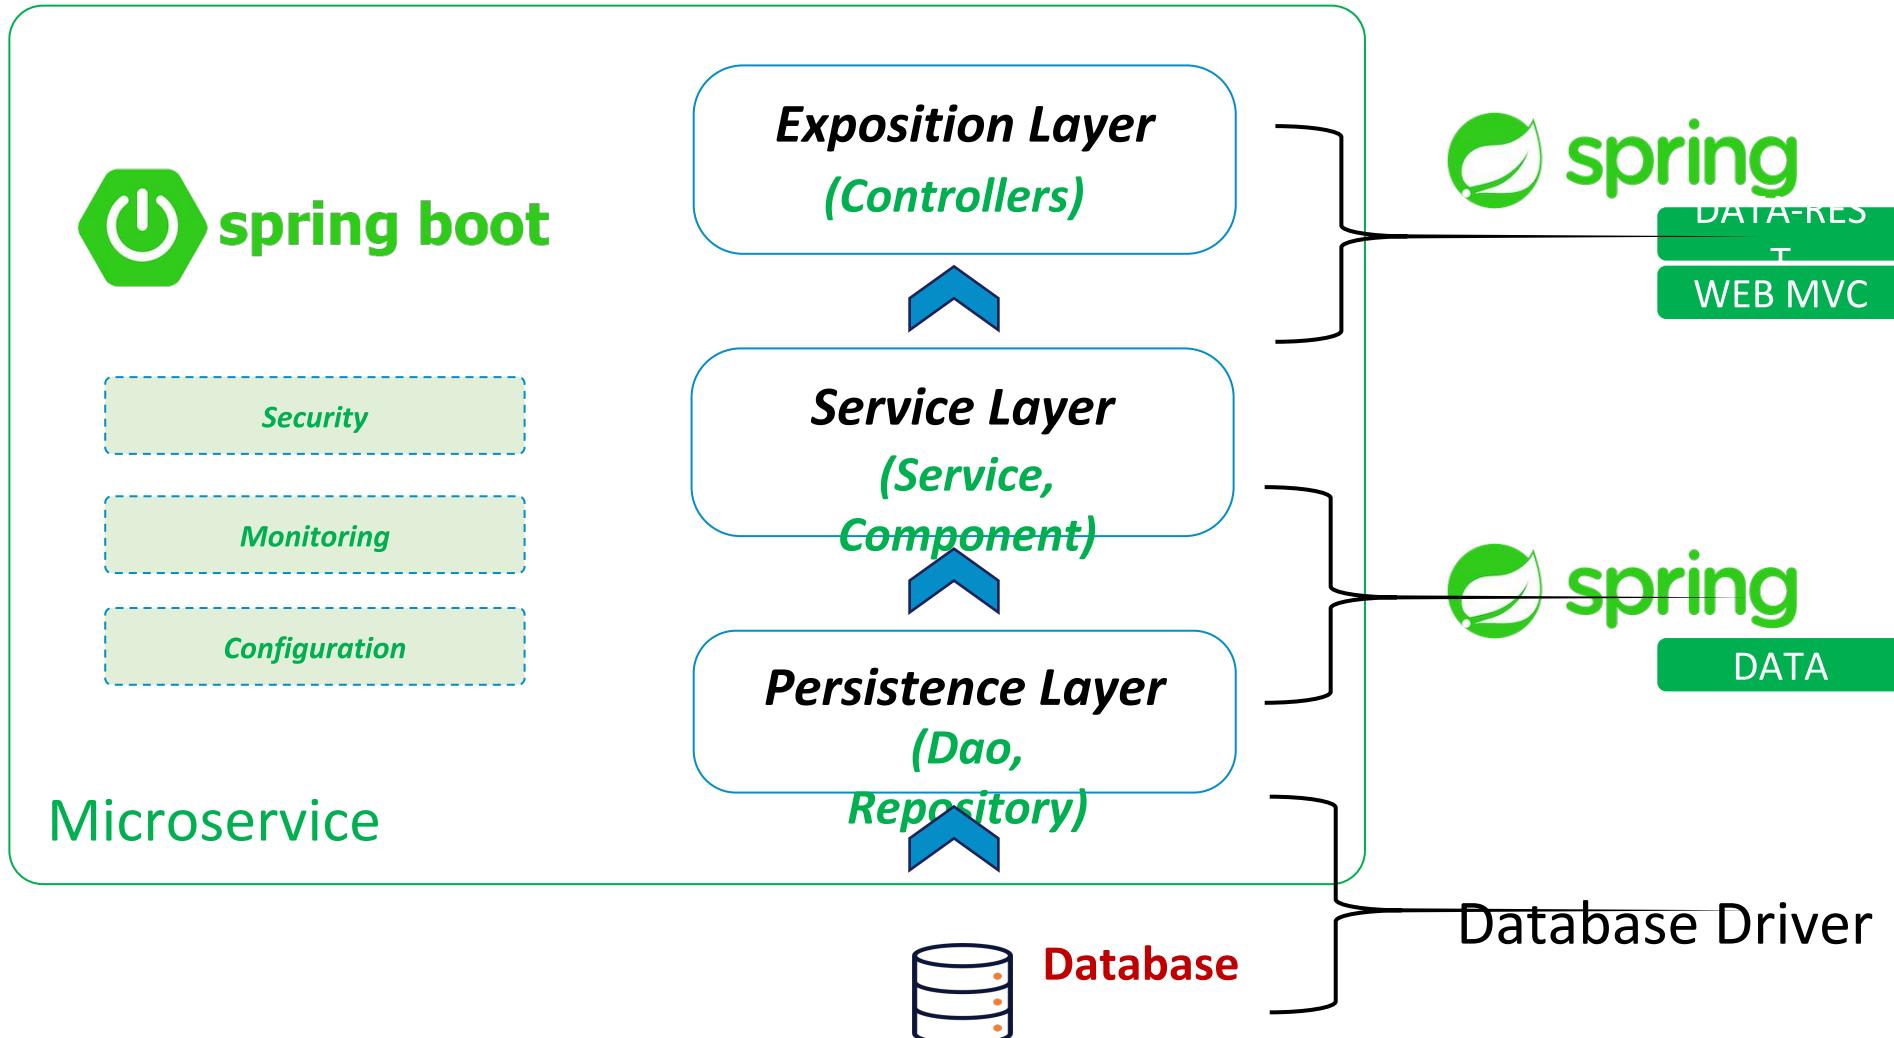
# Exercise 5

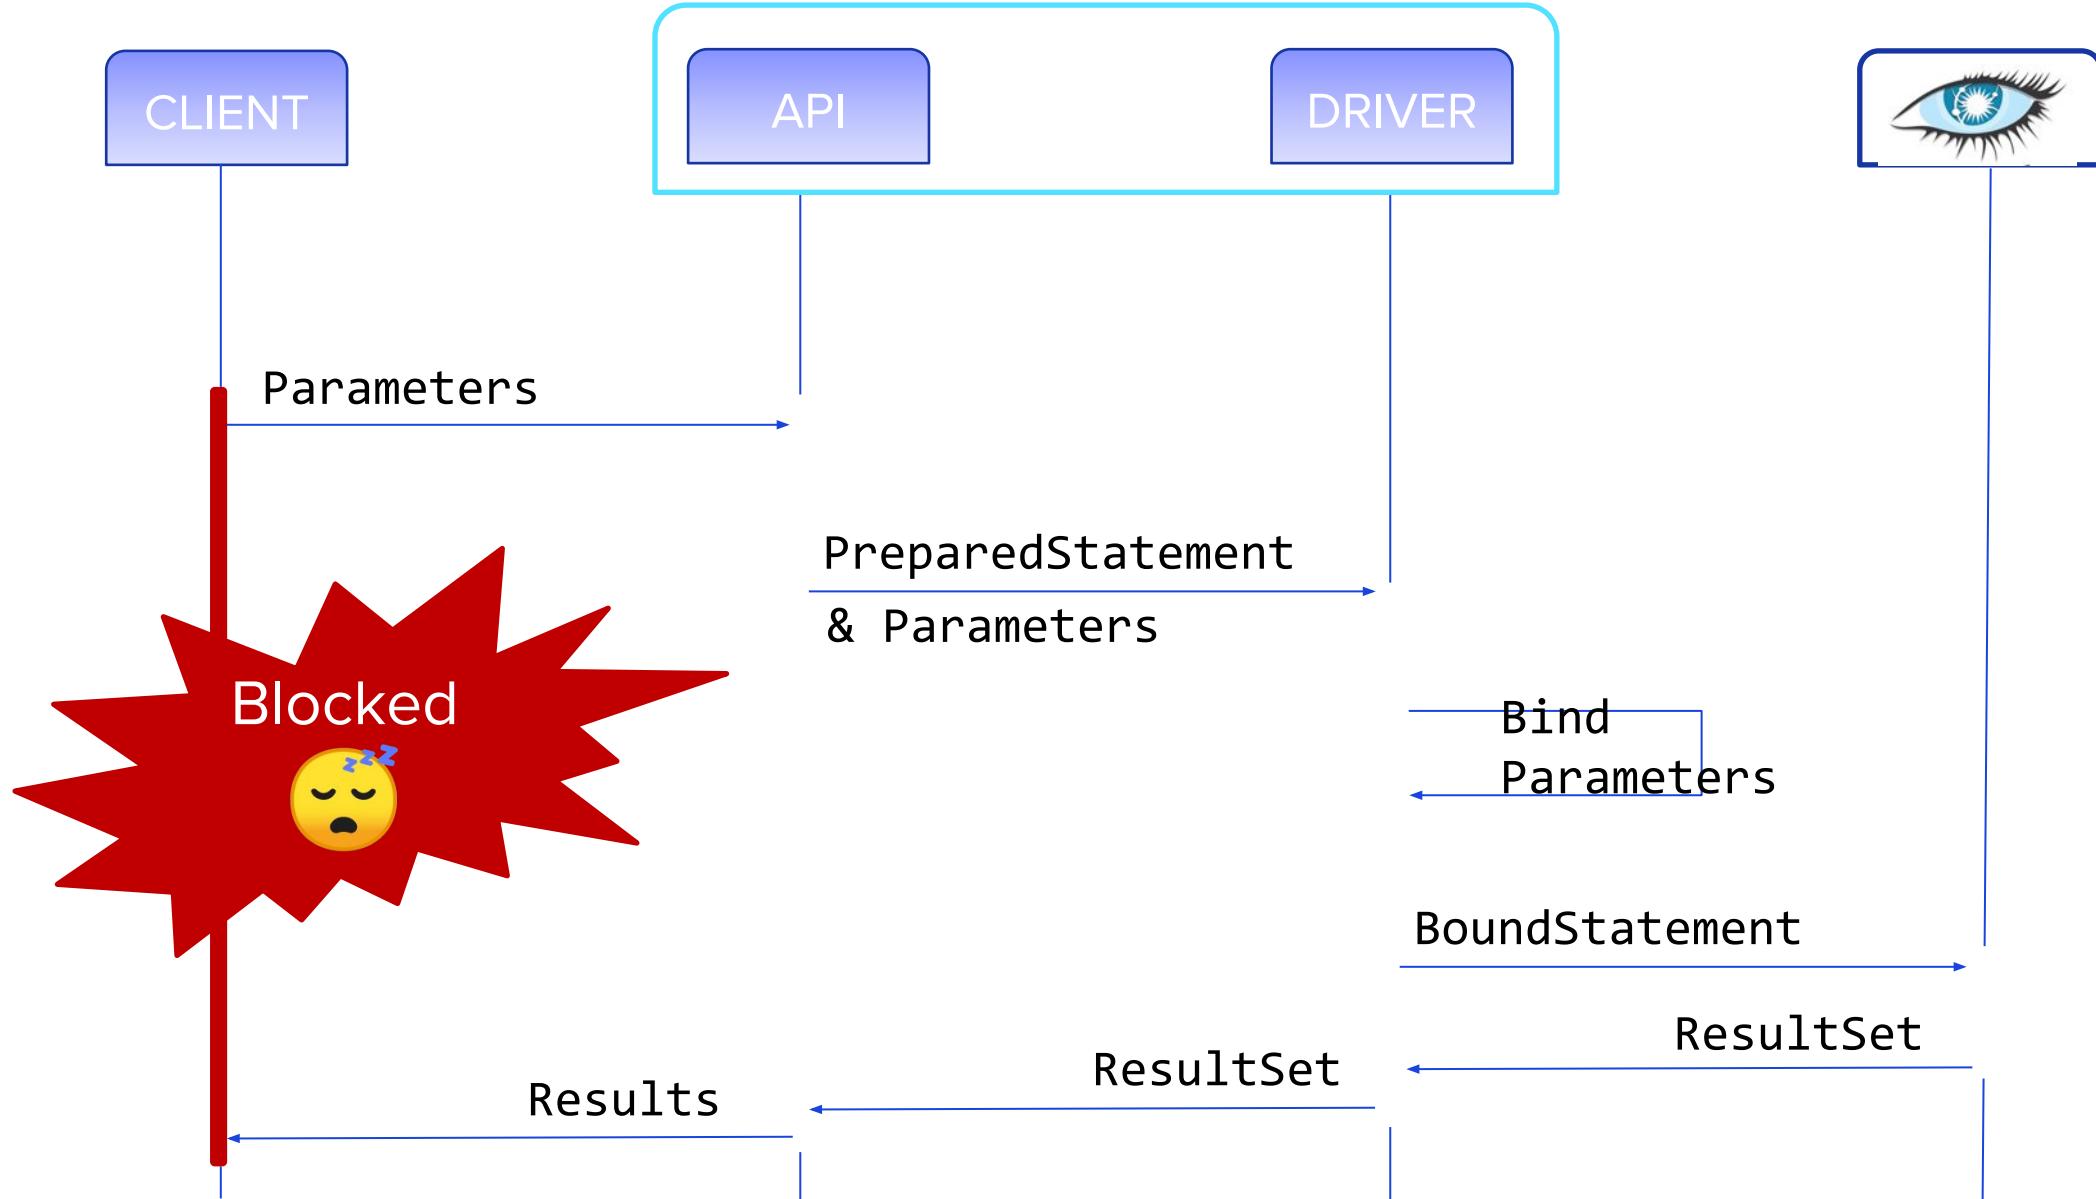
---

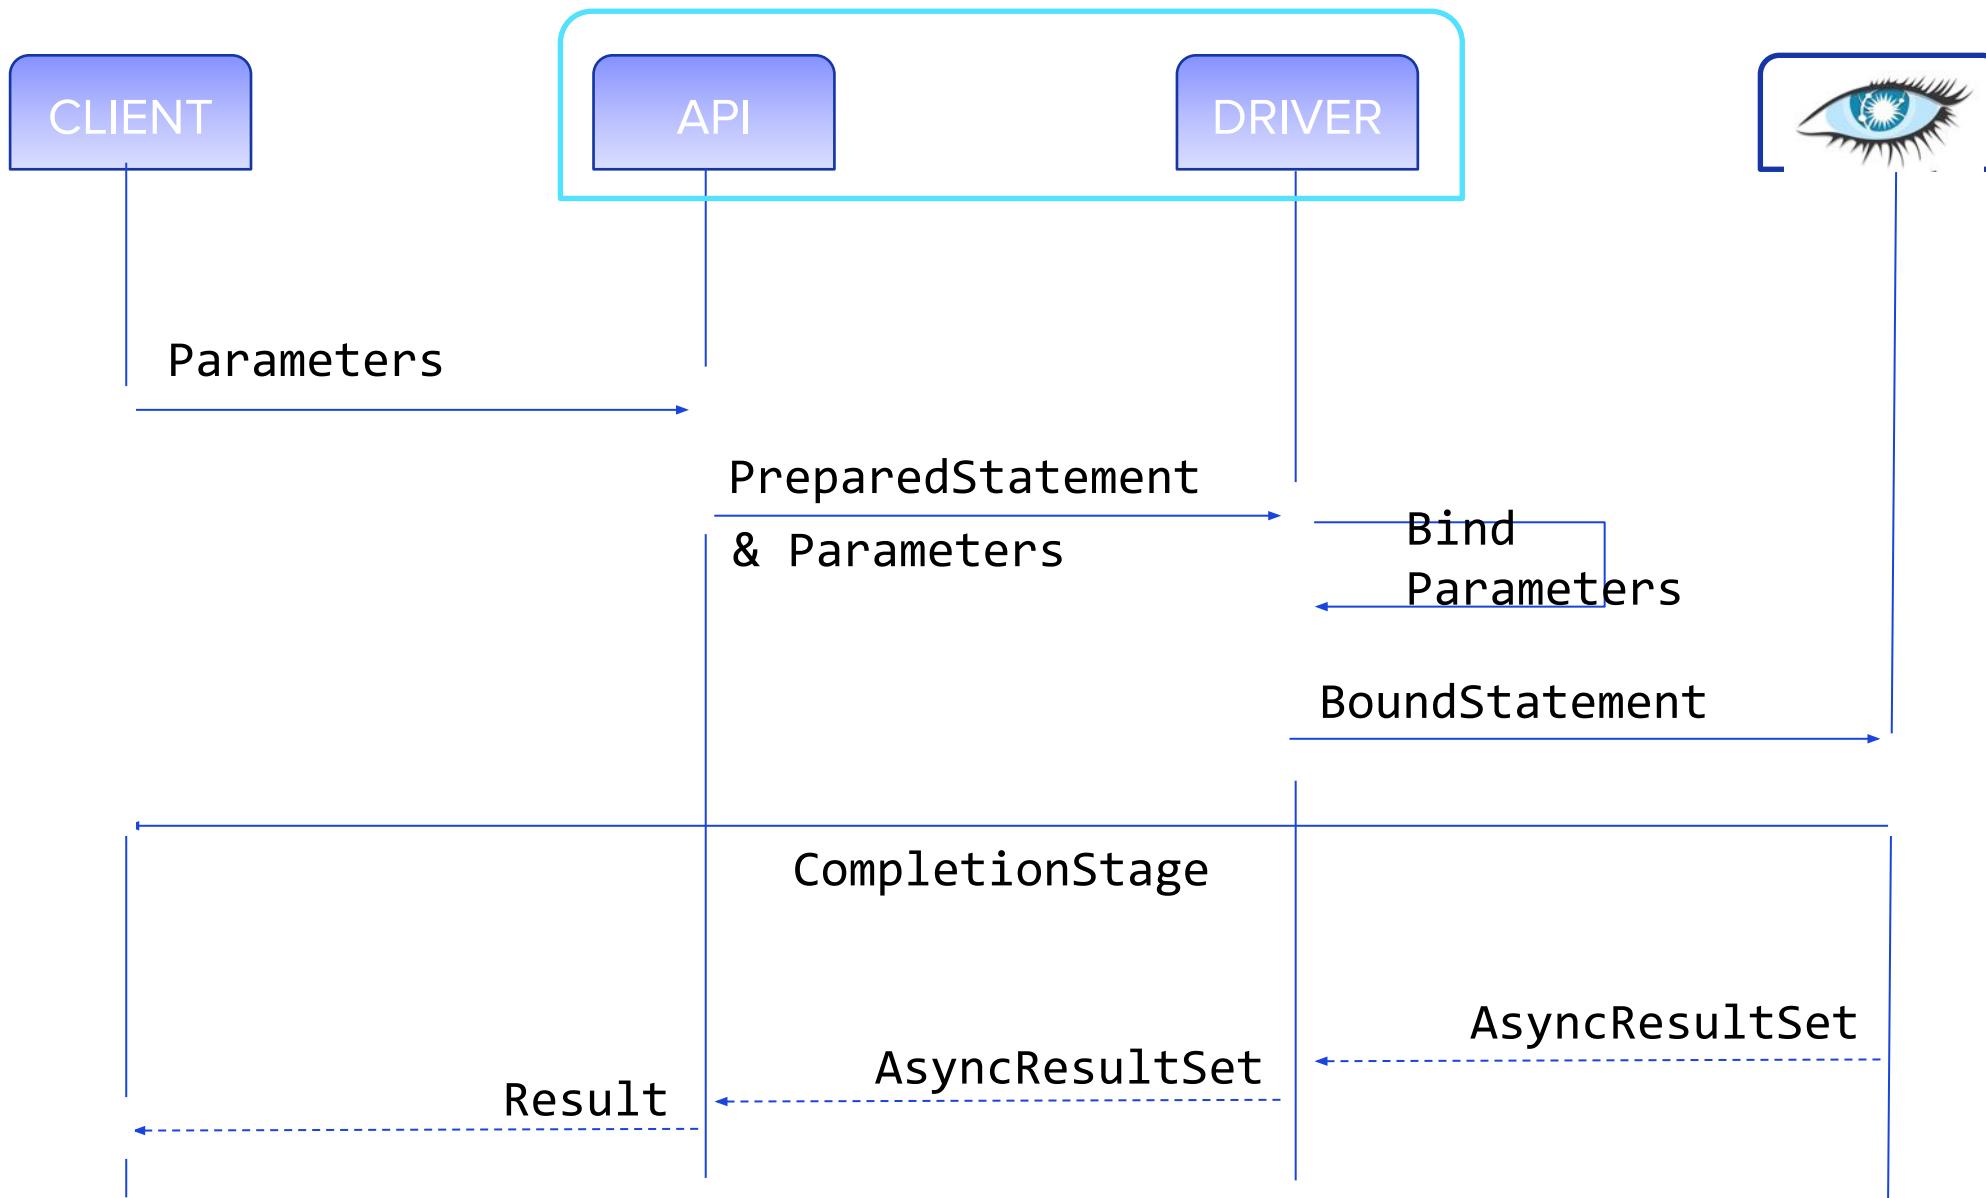


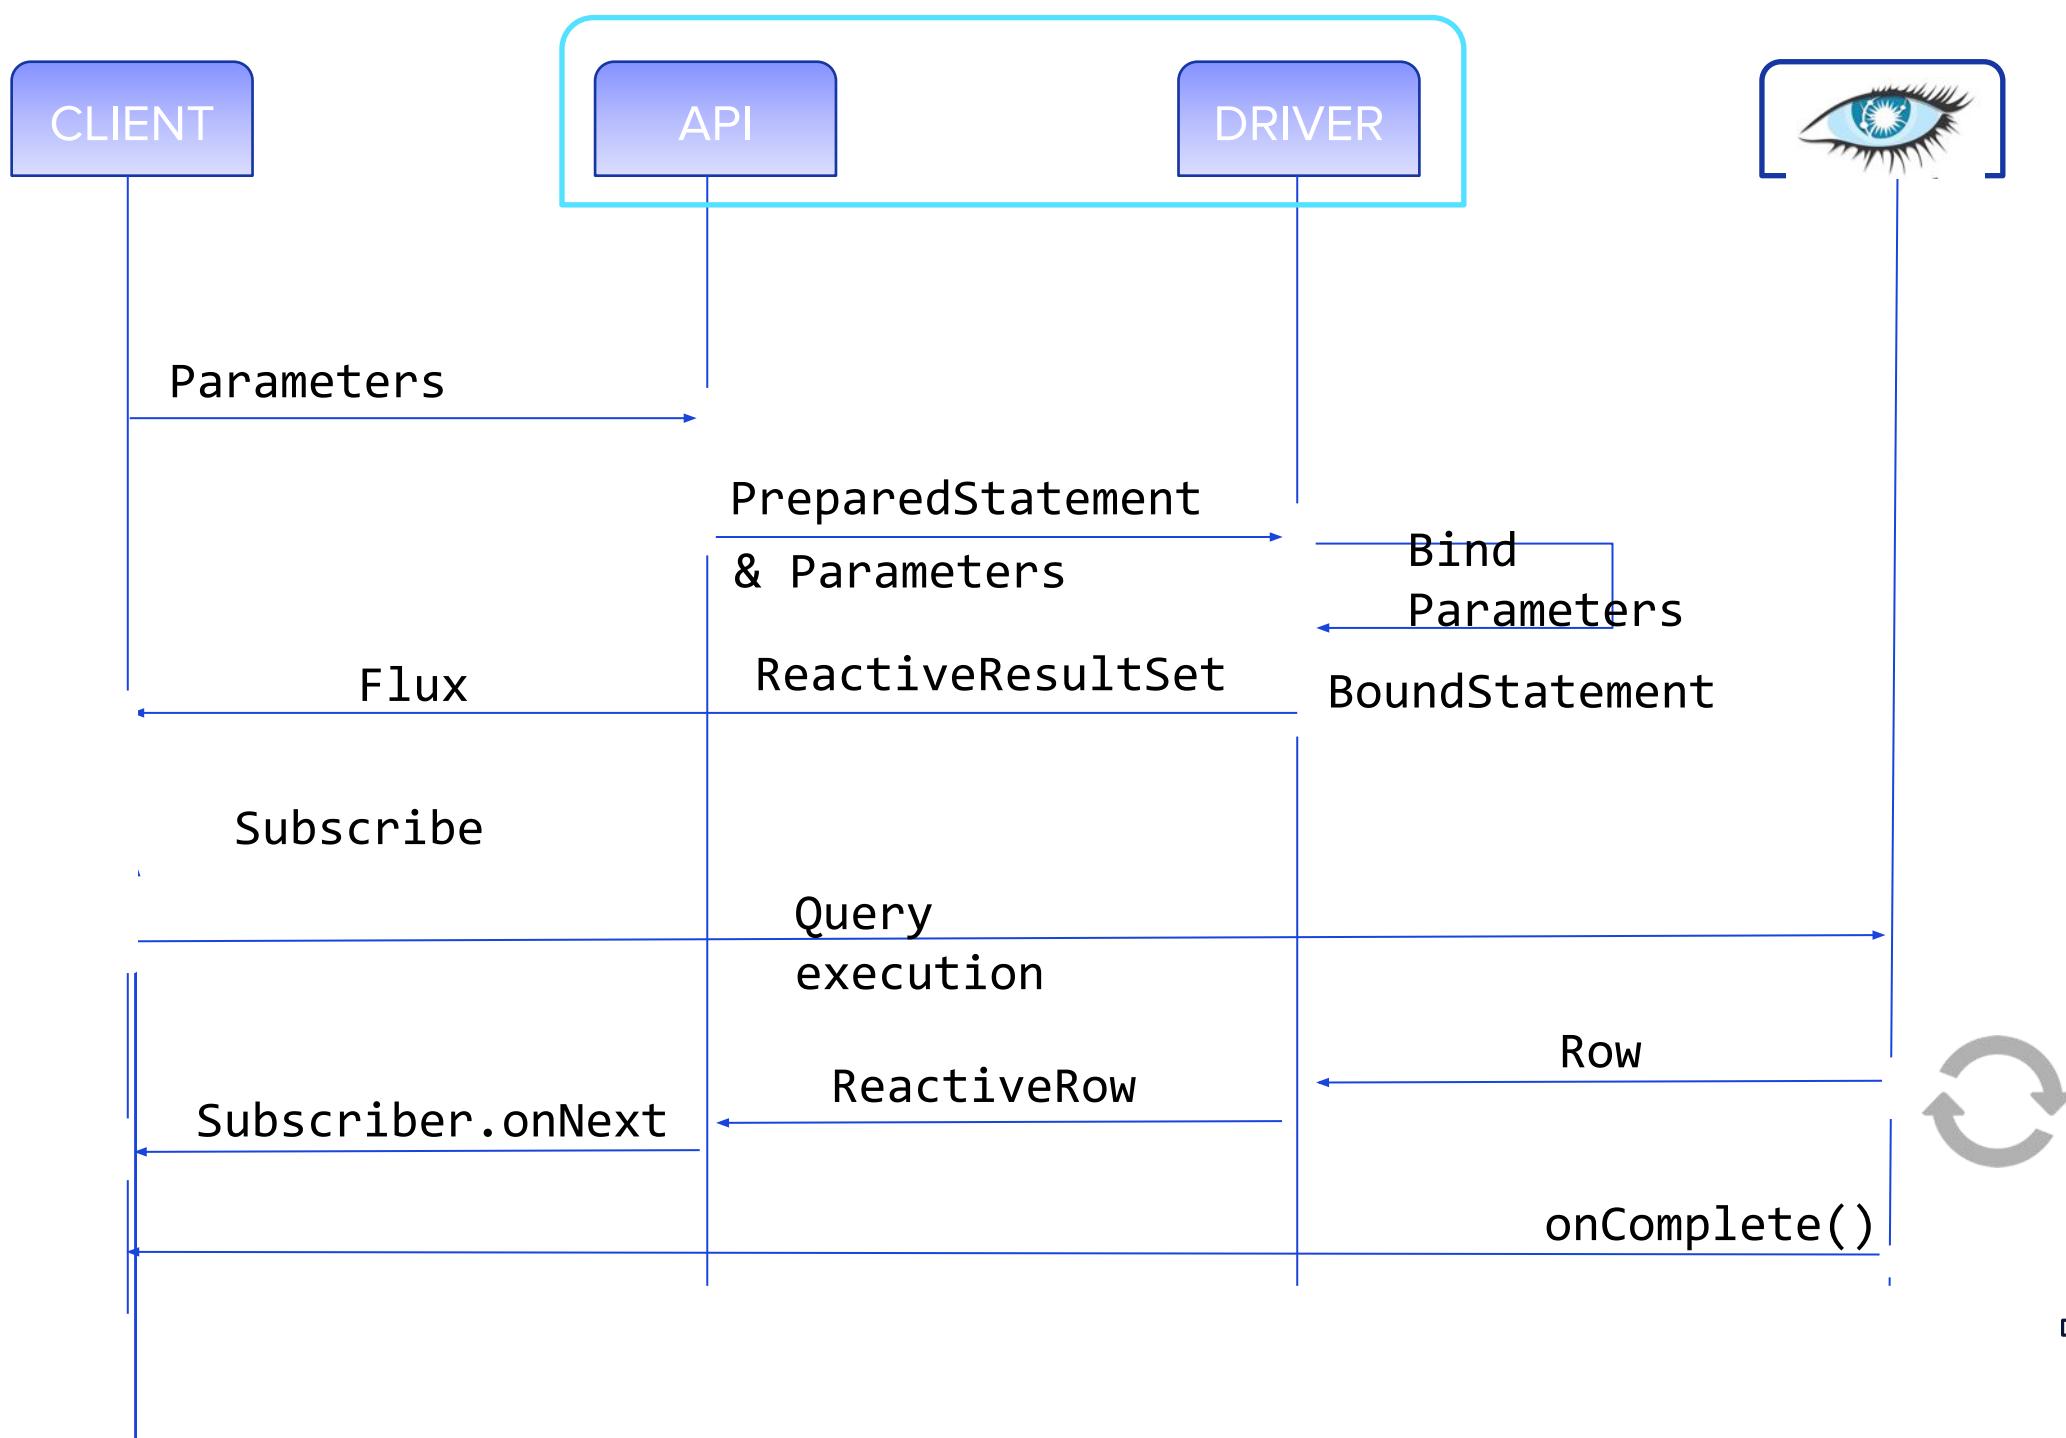
- **Connection with Spring Data**
- **CRUD Repository with Spring Data**
- **Spring Data REST**

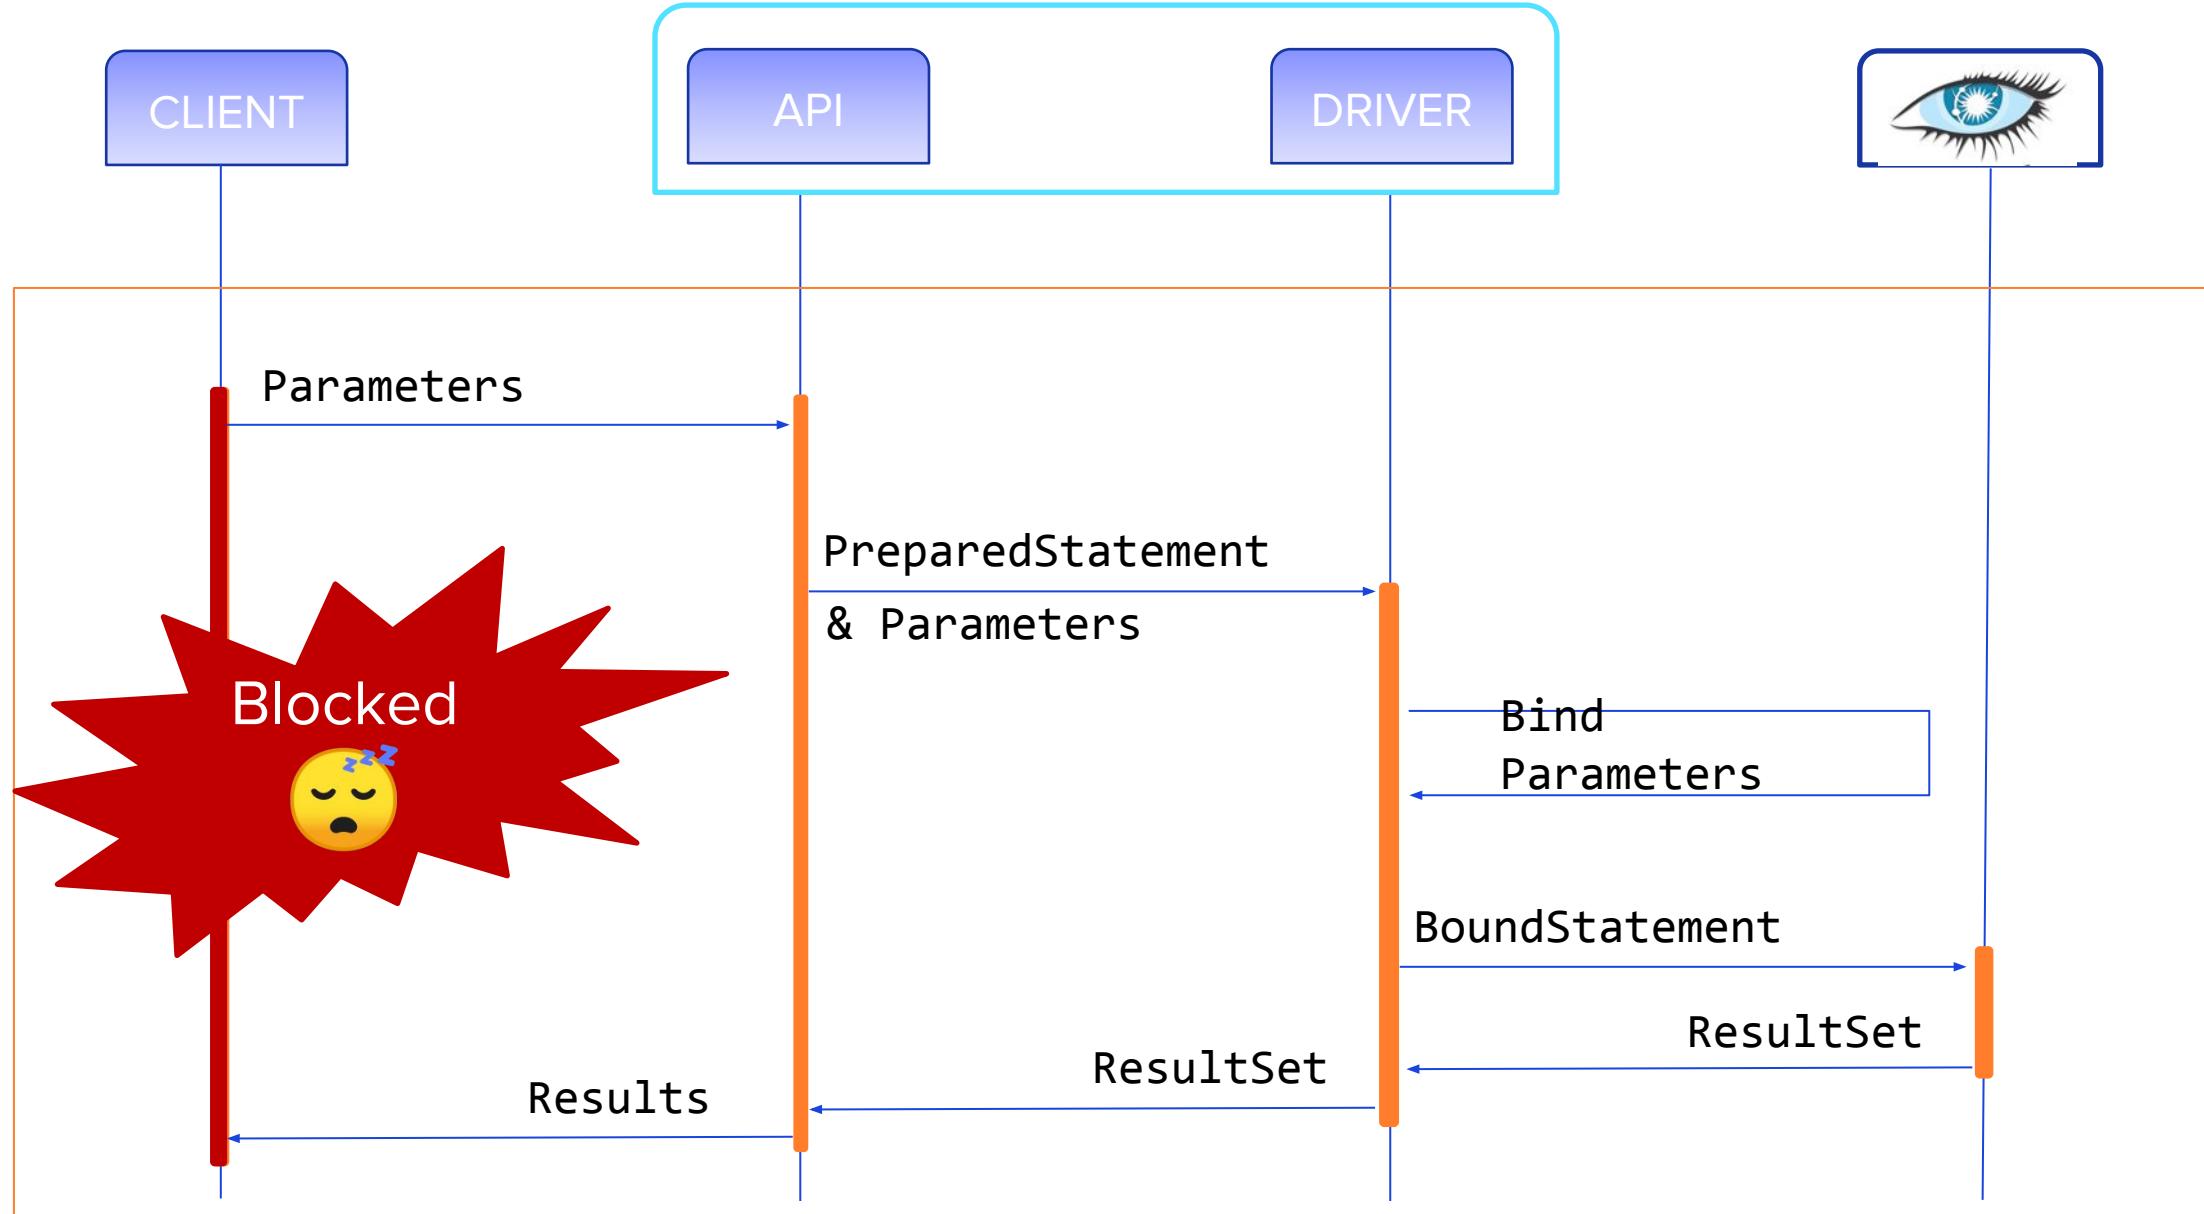


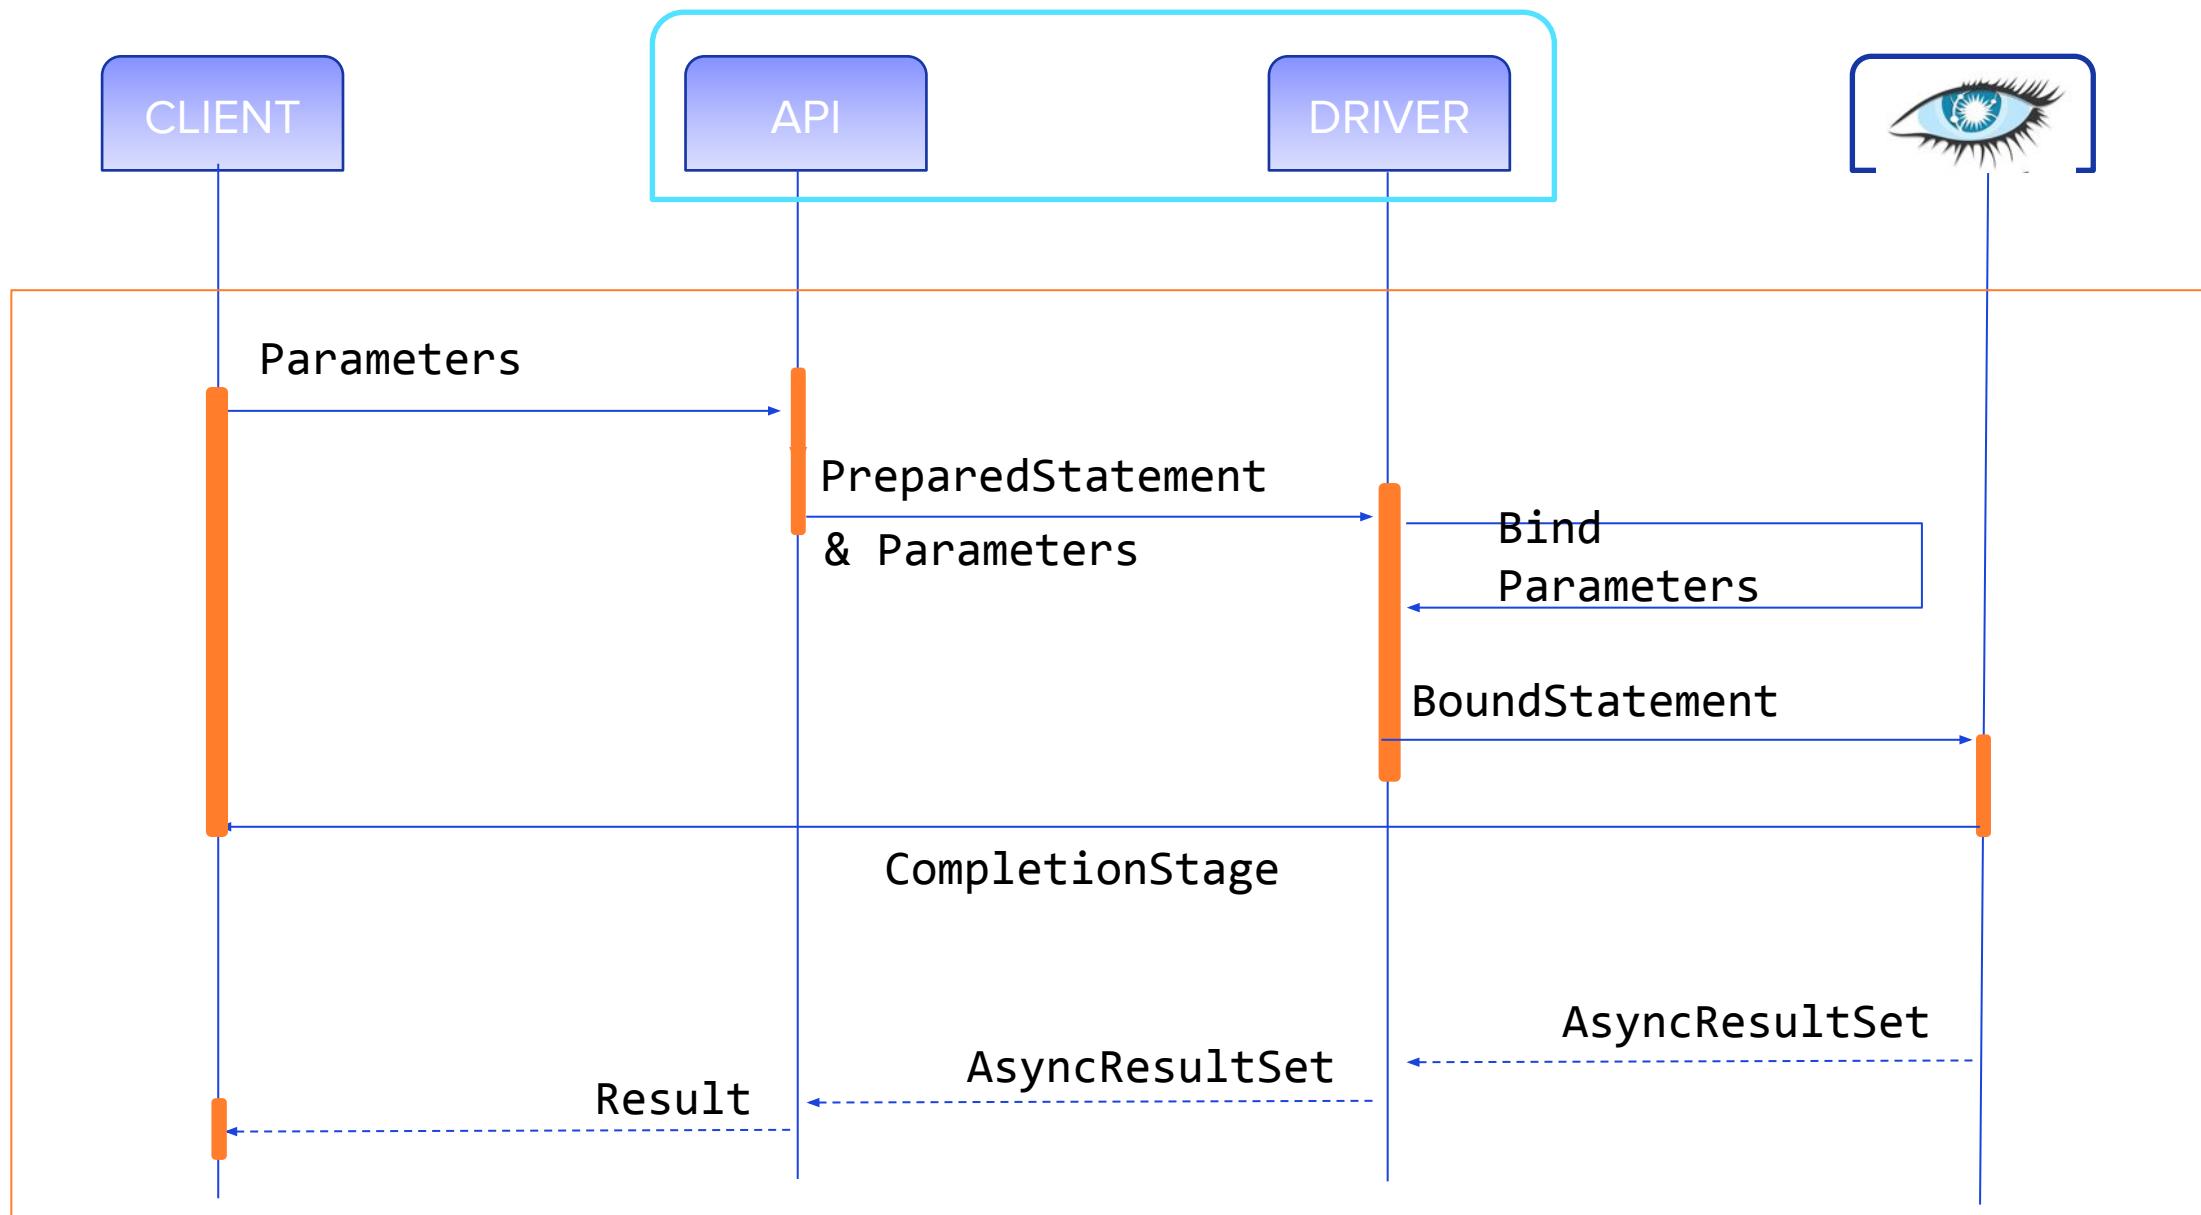












@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>

## Asynchronous Queries



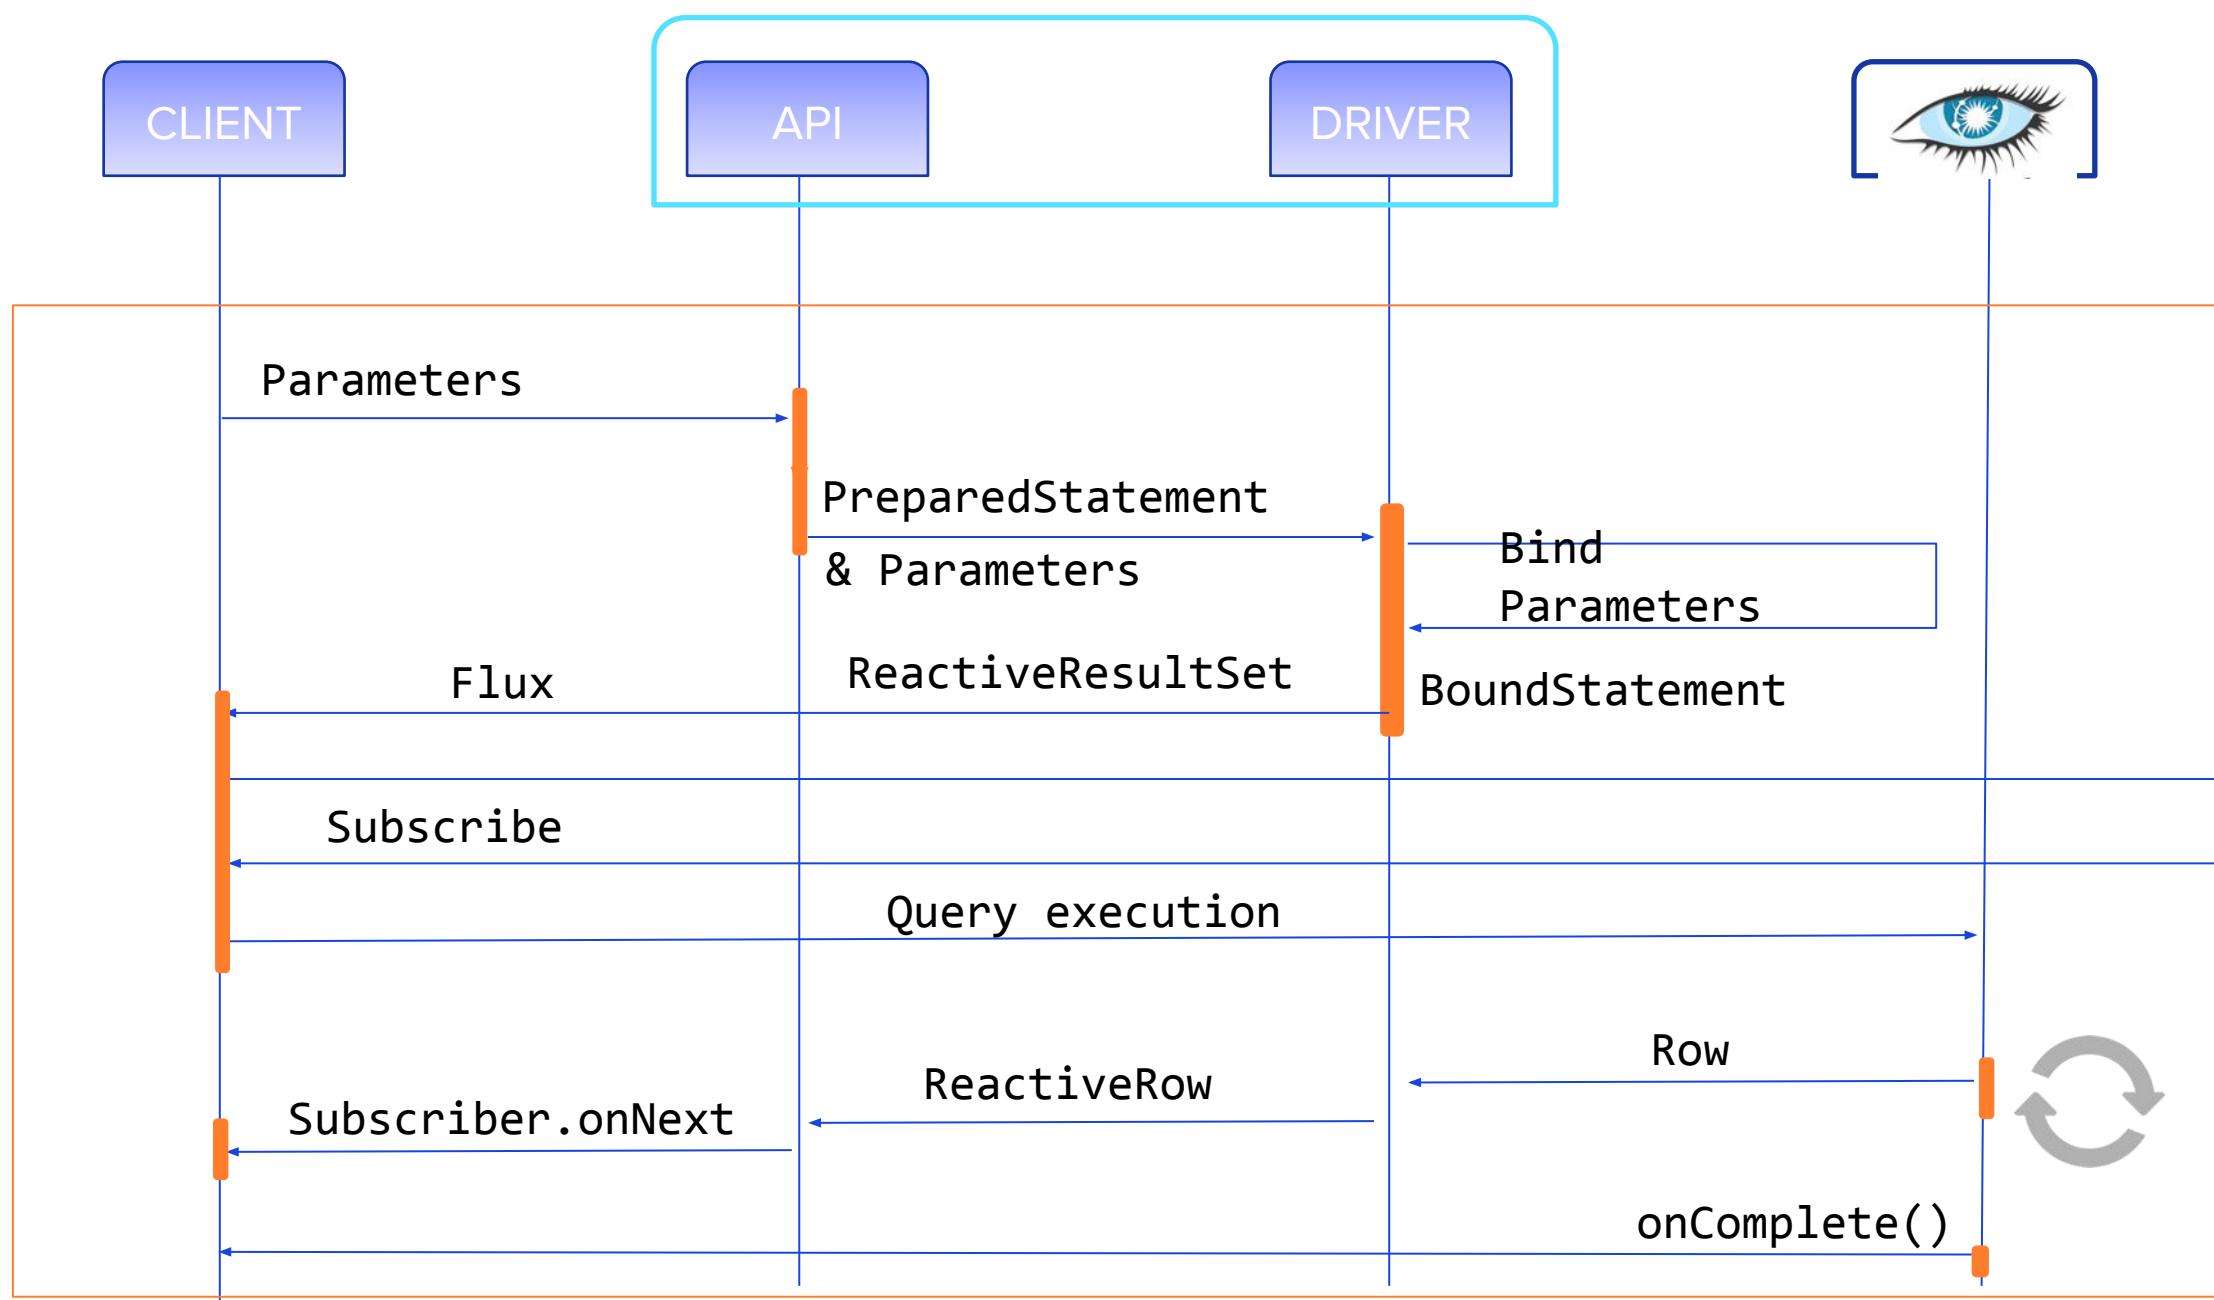
# Asynchronous Queries

```
// From Synchronous
ResultSet resSync = cqlSession.execute(myStatement);

// to Asynchronous
CompletionStage<AsyncResultSet> resAsync =
    cqlSession.executeAsync(myStatement);

resAsync.thenApply(AsyncPagingIterable::one)
    .thenApply(Optional::ofNullable)
    .thenApply(optional -> optional.map(rowMapper))
    .then..
```





# Reactive Queries

```
private final Function<Row, MyBean> rowMapper = ...;

// Execute in reactive way
ReactiveResultSet rs = session.executeReactive(myStatement);

// Return Flux for lists
Flux<ReactiveRow> flux = Flux.from(rs);
Flux<MyBean> res1= flux.skip(offset).take(limit).map(rowMapper);

// Return Mono for single bean
Mono<MyBean> res2 = Mono.fromDirect(rs).map(rowMapper);
```

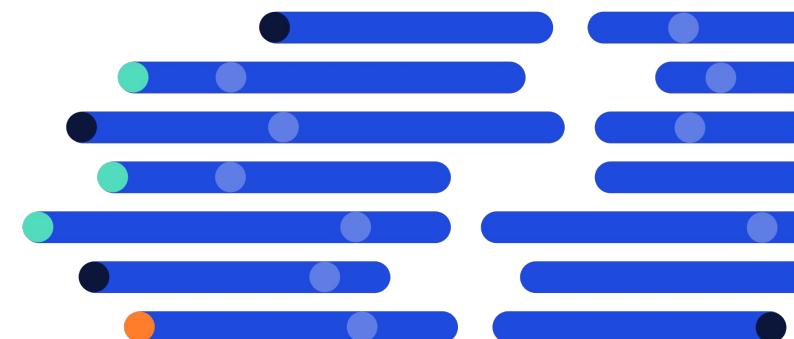


# Exercise 6

---



- **Reactive Repository and Spring WebFlux**
- **Spring Data Reactive**



# Spring Microservices with Apache Cassandra™

- 1 Housekeeping
- 2 Running the TodoApp
- 3 Connectivity to Cassandra
- 4 Create the Repository (dao)
- 5 Spring-Data & Spring-Data-Rest
- 6 Resources

# DataStax Developers on Youtube

- <https://www.youtube.com/channel/UCAIQY251avaMv7bBv5PCo-A>



The screenshot shows the DataStax Developers YouTube channel homepage. At the top, the channel's logo and name "DataStax Developers" are displayed, along with a subscriber count of 2.39K and a "SUBSCRIBED" button. Below the header, there are tabs for "HOME", "VIDEOS" (which is currently selected), "PLAYLISTS", "COMMUNITY", "CHANNELS", and "ABOUT". A search bar and a "SORT BY" dropdown are also present. The main content area displays a grid of video thumbnails. The first two videos in the top row are titled "Cassandra Day Russia" and show a person speaking. The third video in the top row is titled "Distributed Data Show Live Episode 150: Is self driving...". The bottom row includes videos like "Graph Data & Code w/ Denise Gosnell and David Gilardi" and "Building Applications on Astra with BetterBotz...". Each video thumbnail includes its title, duration, and view count.

Video Title	Duration	Views
Cassandra Day Russia [CL Room] Воркшопы	6:25:54	551 views • 16 hours ago
Cassandra Day Russia [RF Room] Доклады	4:04:51	192 views • 16 hours ago
Distributed Data Show Live Episode 150: Is self driving...	43:25	30 views • 18 hours ago
Testing Scalable Distributed Systems: Workshop with...	2:34:15	297 views • 2 days ago
Graph Data & Code w/ Denise Gosnell and David Gilardi	55:46	69 views • 3 days ago
Cassandra Developer Workshop GMT	3:43:33	3.6K views • 3 days ago
What do you need to know to deploy Cassandra in...	21:10	94 views • 1 week ago
Building Applications on Astra with BetterBotz...	1:59:26	1.3K views • 1 week ago

# Sample CODES



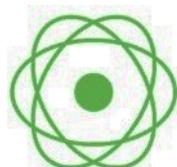
Micro-service **REST**

<https://bit.ly/31RL62I>



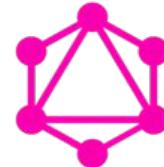
Micro-service **GRPC**

<https://bit.ly/2pofk0b>



**Reactive with Webflux**

<https://bit.ly/34ePzhL>



Micro-service **GraphQL**

<https://bit.ly/2MVicup>



Micro-service **Kafka**

<https://bit.ly/2JwsFdM>



Micro-service **Serverless**

<https://bit.ly/31VQz8G>

# menti.com

00 00 00



Available on the iPhone  
**App Store**

GET IT ON  
**Google play**

# Developer Resources

**LEARN**

Join [academy.datastax.com](https://academy.datastax.com)  
Free online courses - Cassandra certifications

**ASK/SERVE**

Join [community.datastax.com](https://community.datastax.com)  
Ask/answer community user questions - share your expertise

 [Subscribe](#)

 [You Tube](#)



<https://bit.ly/cassandra-cert-FREE>

# Get Certified !



	Administrator	Architect	Developer	Graph Specialist
DS101: Introduction to Apache Cassandra™	Included			
DS201: DataStax Enterprise 6 Foundations of Apache Cassandra™	Included			
DS210: DataStax Enterprise 6 Operations with Apache Cassandra™	Included			
DS220: DataStax Enterprise 6 Practical Application Data Modeling with Apache Cassandra™	N/A			
DS330: DataStax Enterprise 6 Graph	N/A			

# Cassandra The Definitive Guide

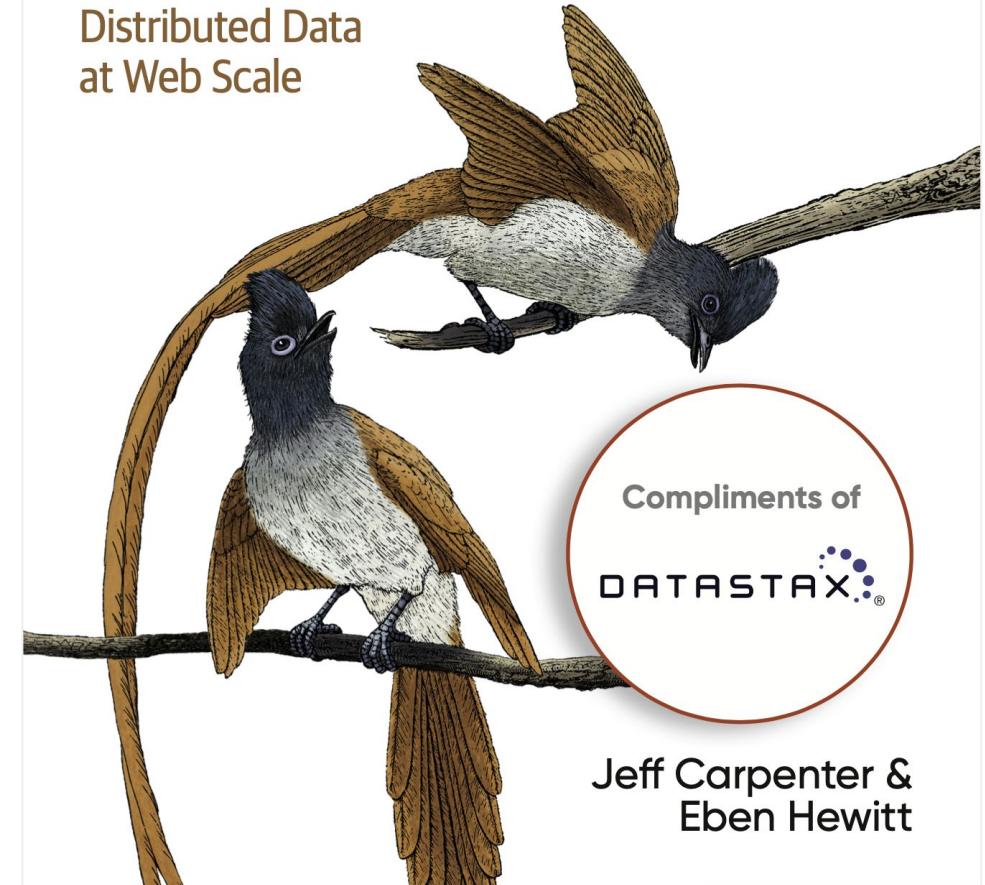
<https://www.datastax.com/resources/ebook/oreilly-cassandra-definitive-guide>

O'REILLY®

Third  
Edition

# Cassandra The Definitive Guide

Distributed Data  
at Web Scale



Jeff Carpenter &  
Eben Hewitt

# Upcoming events

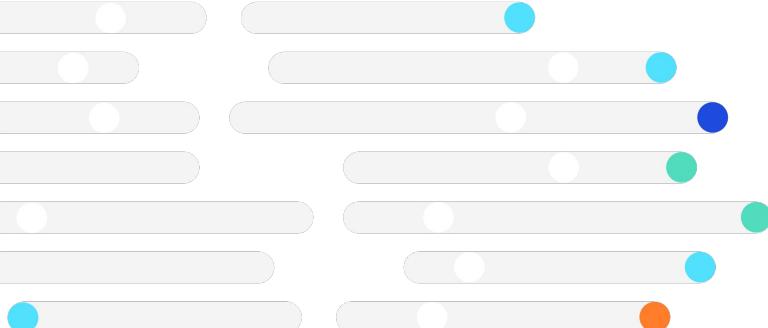
Date	Time	Content	Type
6/24	12pm IST	Build a REST API for Cassandra with NodeJS and Python	WORKSHOP = HANDS-ON
July 1st July 2nd	12pm EDT 12:30pm IST		
July 8th July 9th			
7/15			
7/22			
7/29			



@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>





---

**Thank you**